



SELA|DEVELOPER|PRACTICE
December 6-8, 2020

Leon Jalfon

Service Mesh Fundamentals

Agenda

- ❖ Workshop Introduction
- ❖ Understanding Service Mesh
- ❖ Service Mesh Fundamentals with Linkerd
- ❖ Service Mesh Fundamentals with Istio
- ❖ Service Mesh Interface (SMI)
- ❖ Summary

Workshop Objectives

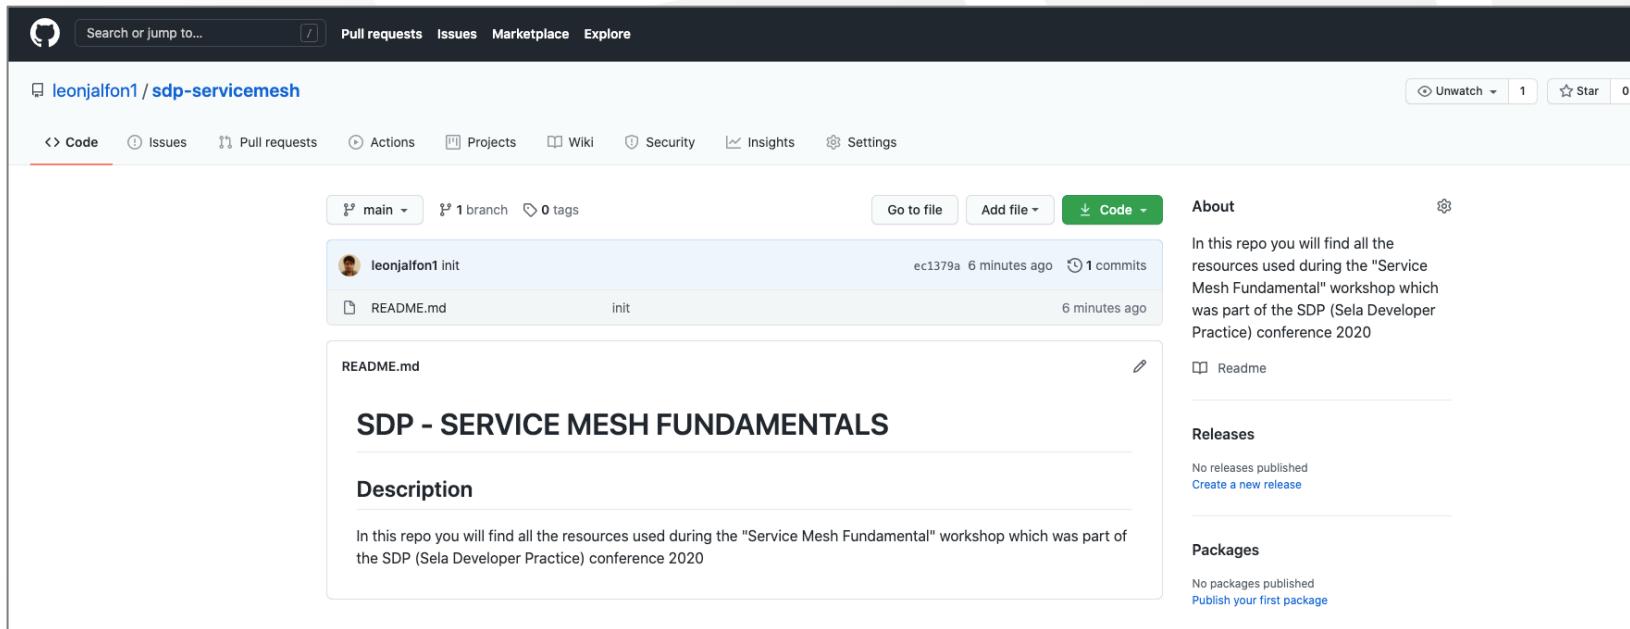
- ★ Getting started with service mesh on Kubernetes
- ★ Understand how service mesh works
- ★ Learn Linkerd and Istio through hands-on labs
- ★ And most importantly, let's have fun!

Knowledge Assumptions

- ★ Linux knowledge
- ★ Networking understanding
- ★ Hands-on experience with Docker (or another container technology)
- ★ Hands-on experience with Kubernetes (especially regarding networking)
- ★ Experience managing Microservices on Kubernetes
- ★ Kubernetes CRD's and sidecar pattern understanding

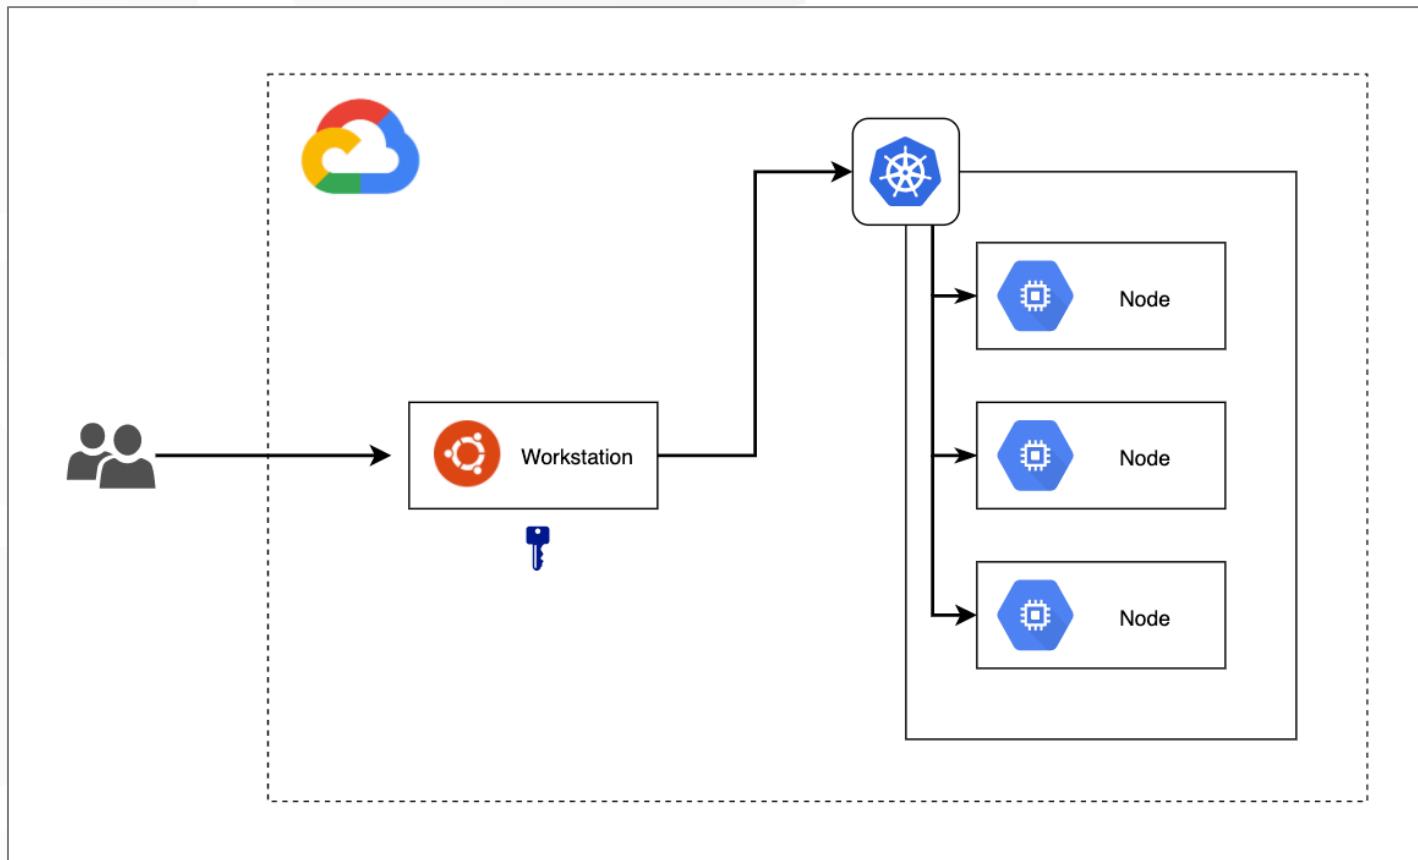
Workshop Resources

- ★ You can find all the resources used during the workshop in my GitHub



<https://github.com/leonjalfon1/sdp-servicemesh>

Lab Infrastructure



Workshop Hand-On Labs - Linkerd

- ★ LAB 01: Installing Linkerd
- ★ LAB 02: Exploring Linkerd
- ★ LAB 03: Debugging with Linkerd
- ★ LAB 04: Validating Linkerd mTLS
- ★ LAB 05: Deploying and Debugging the BookApp
- ★ LAB 06: Working with Service Profiles

Workshop Hand-On Labs - Istio

- ★ LAB 07: Installing Istio
- ★ LAB 08: Exploring Istio
- ★ LAB 09: Intelligent Routing with Istio
- ★ LAB 10: Fault Injection with Istio
- ★ LAB 11: Traffic Management with Istio
- ★ LAB 12: Setup Istio Authorization

LETS GET STARTED





SERVICE MESH



Open
Service
Mesh



SERVICE MESH EVERYWHERE



LINKERD

What is a Service Mesh?

Not logged in [Talk](#) [Contributions](#) [Create account](#) [Log in](#)

Article [Talk](#) Read [Edit](#) [View history](#) [Search Wikipedia](#) 

Service mesh

From Wikipedia, the free encyclopedia



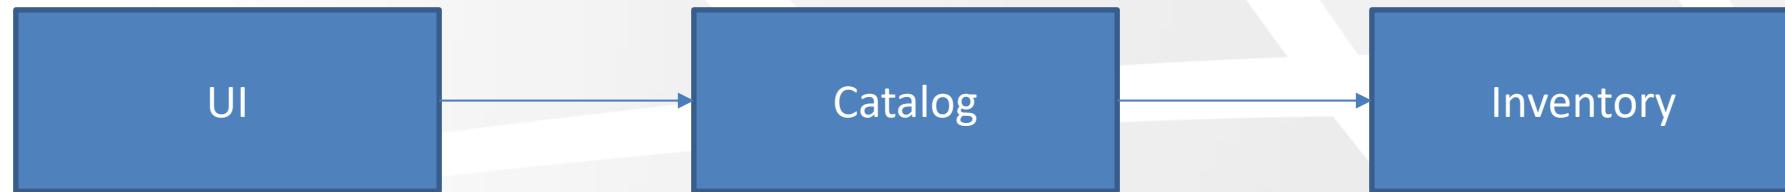
This article has multiple issues. Please help [improve it](#) or [\[show\]](#) discuss these issues on the [talk page](#). (*Learn how and when to remove these template messages*)

In software architecture, a service mesh is a dedicated infrastructure layer for facilitating service-to-service communications between microservices, often using a sidecar proxy.

Having such a dedicated communication layer can provide a number of benefits, such as providing observability into communications, providing secure connections, or automating retries and backoff for failed requests.

Understanding Service Mesh

Common application without using a service mesh



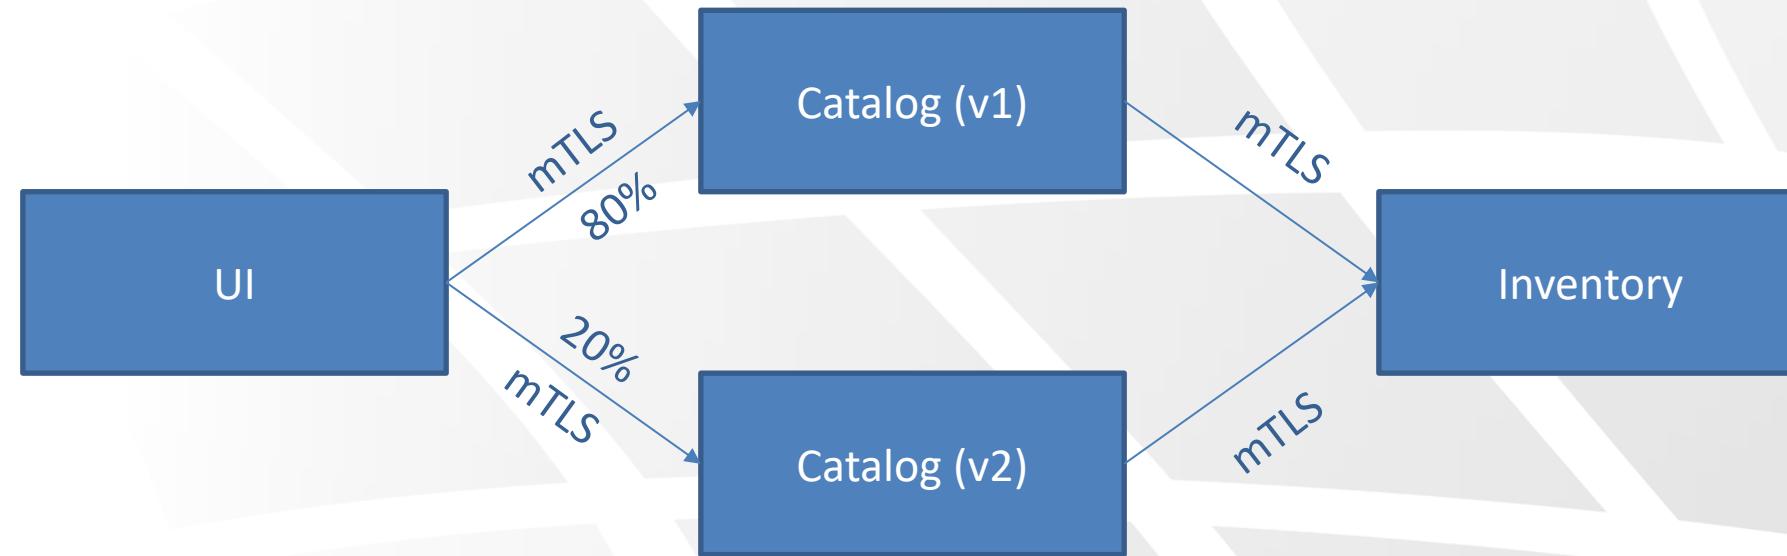
Understanding Service Mesh

Improve Security → Automatic mutual TLS



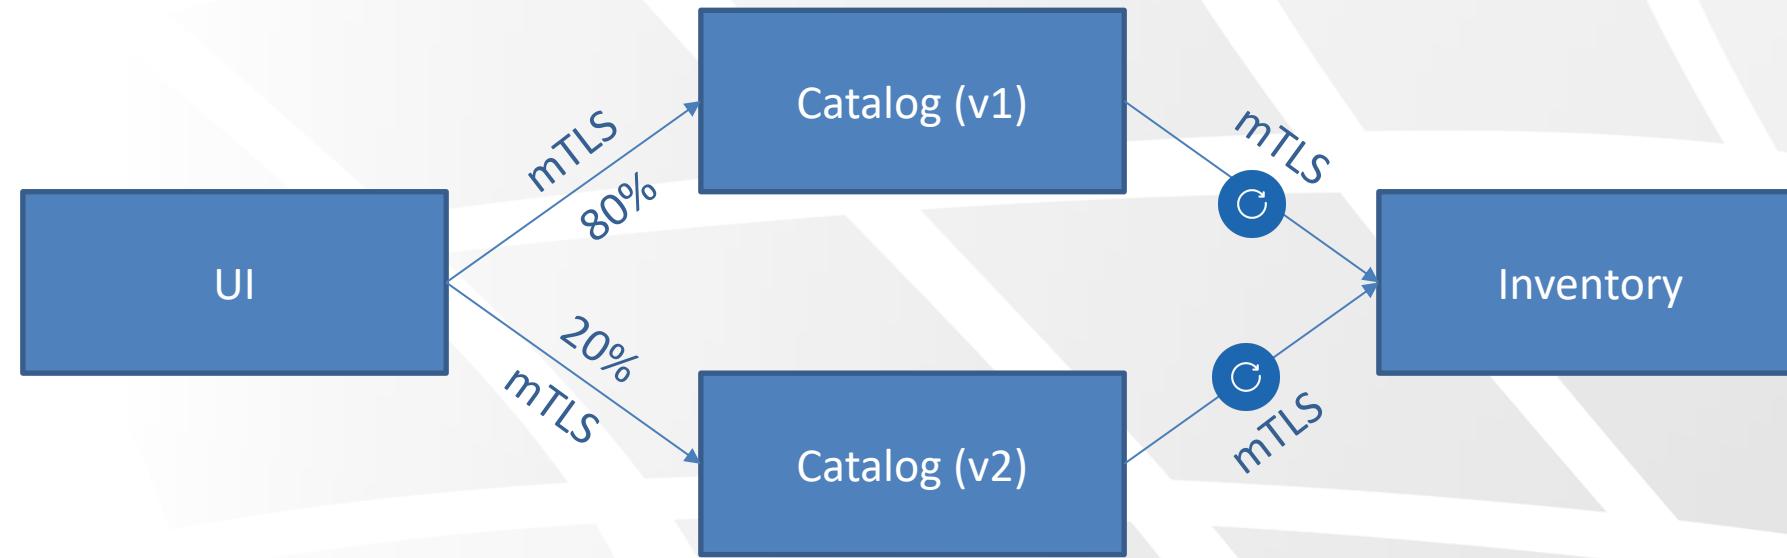
Understanding Service Mesh

Improve Networking → Intelligent Routing (canary deployments, blue/green, etc)



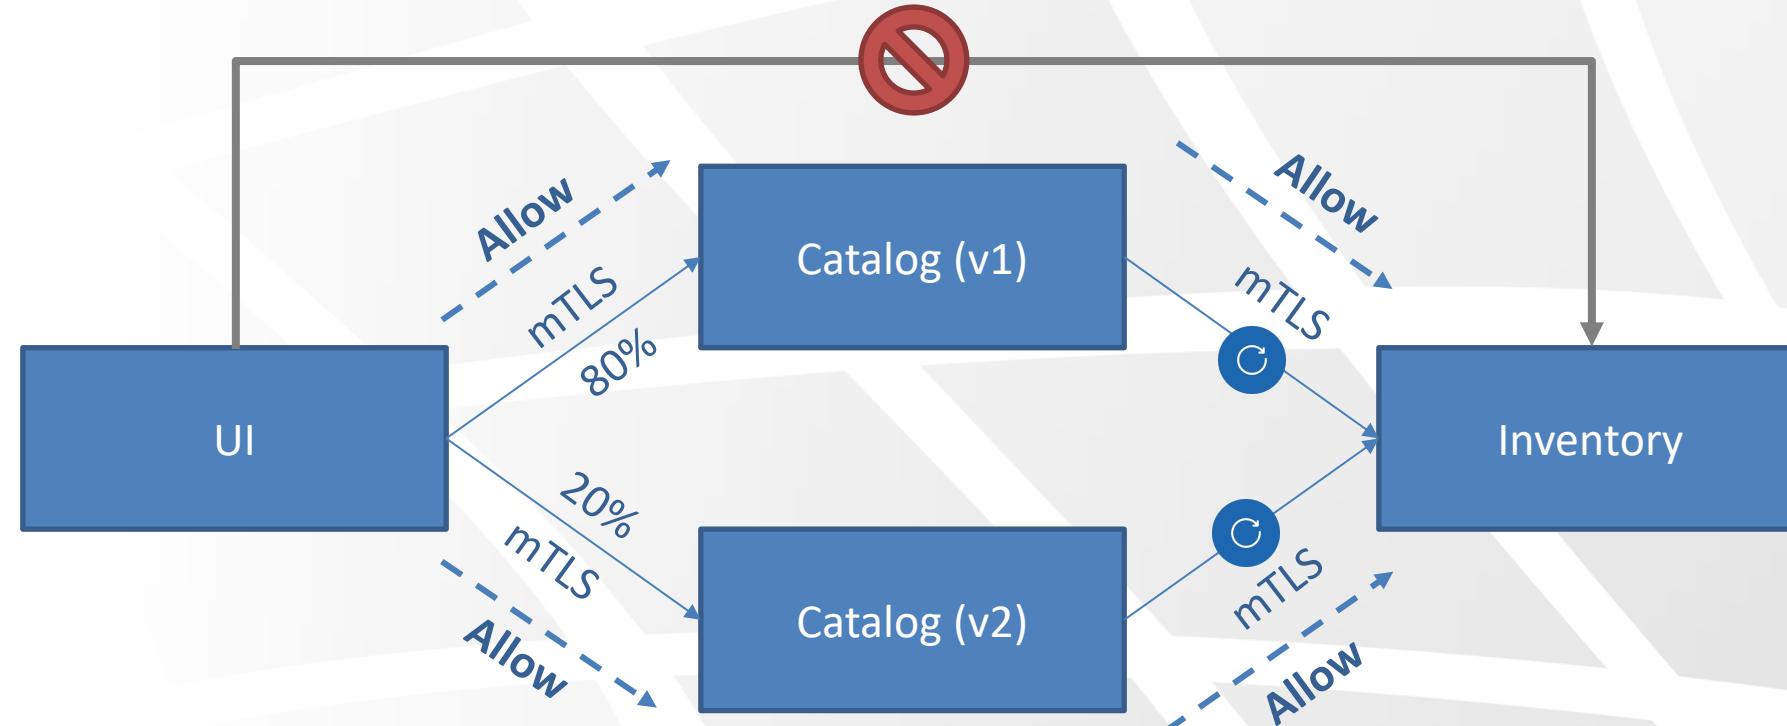
Understanding Service Mesh

Improve Networking → Network Policies (retries, timeouts, fault injections, etc)



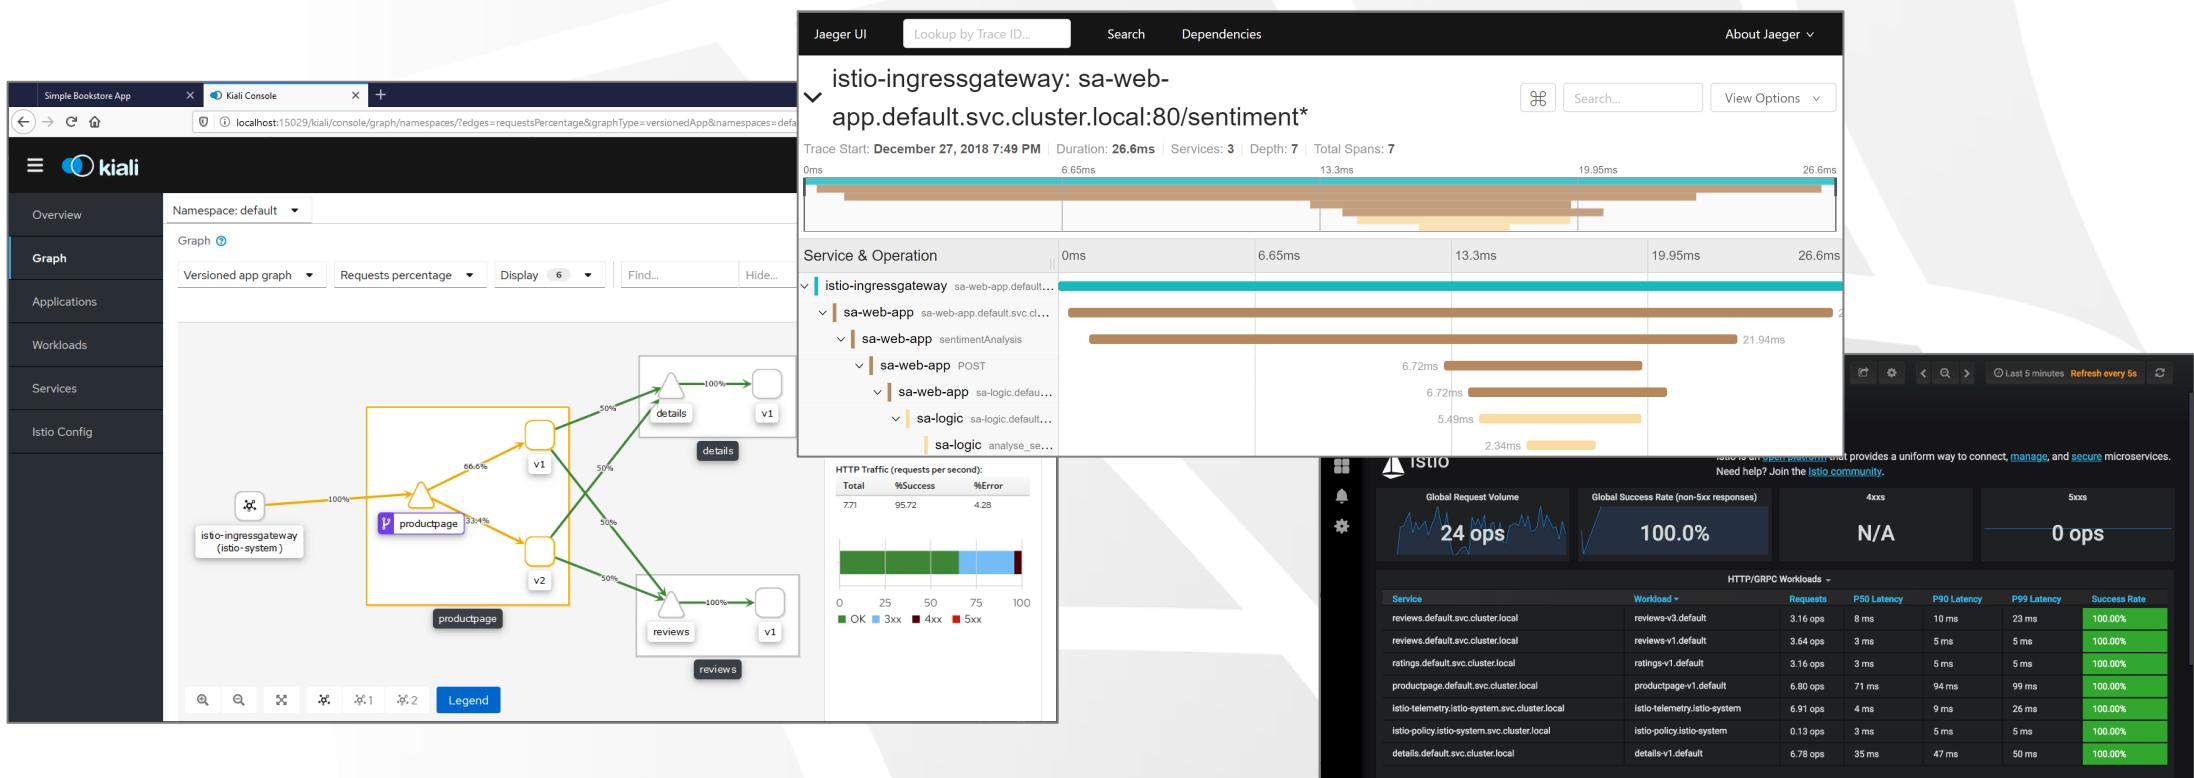
Understanding Service Mesh

Improve Control → Authorization Rules

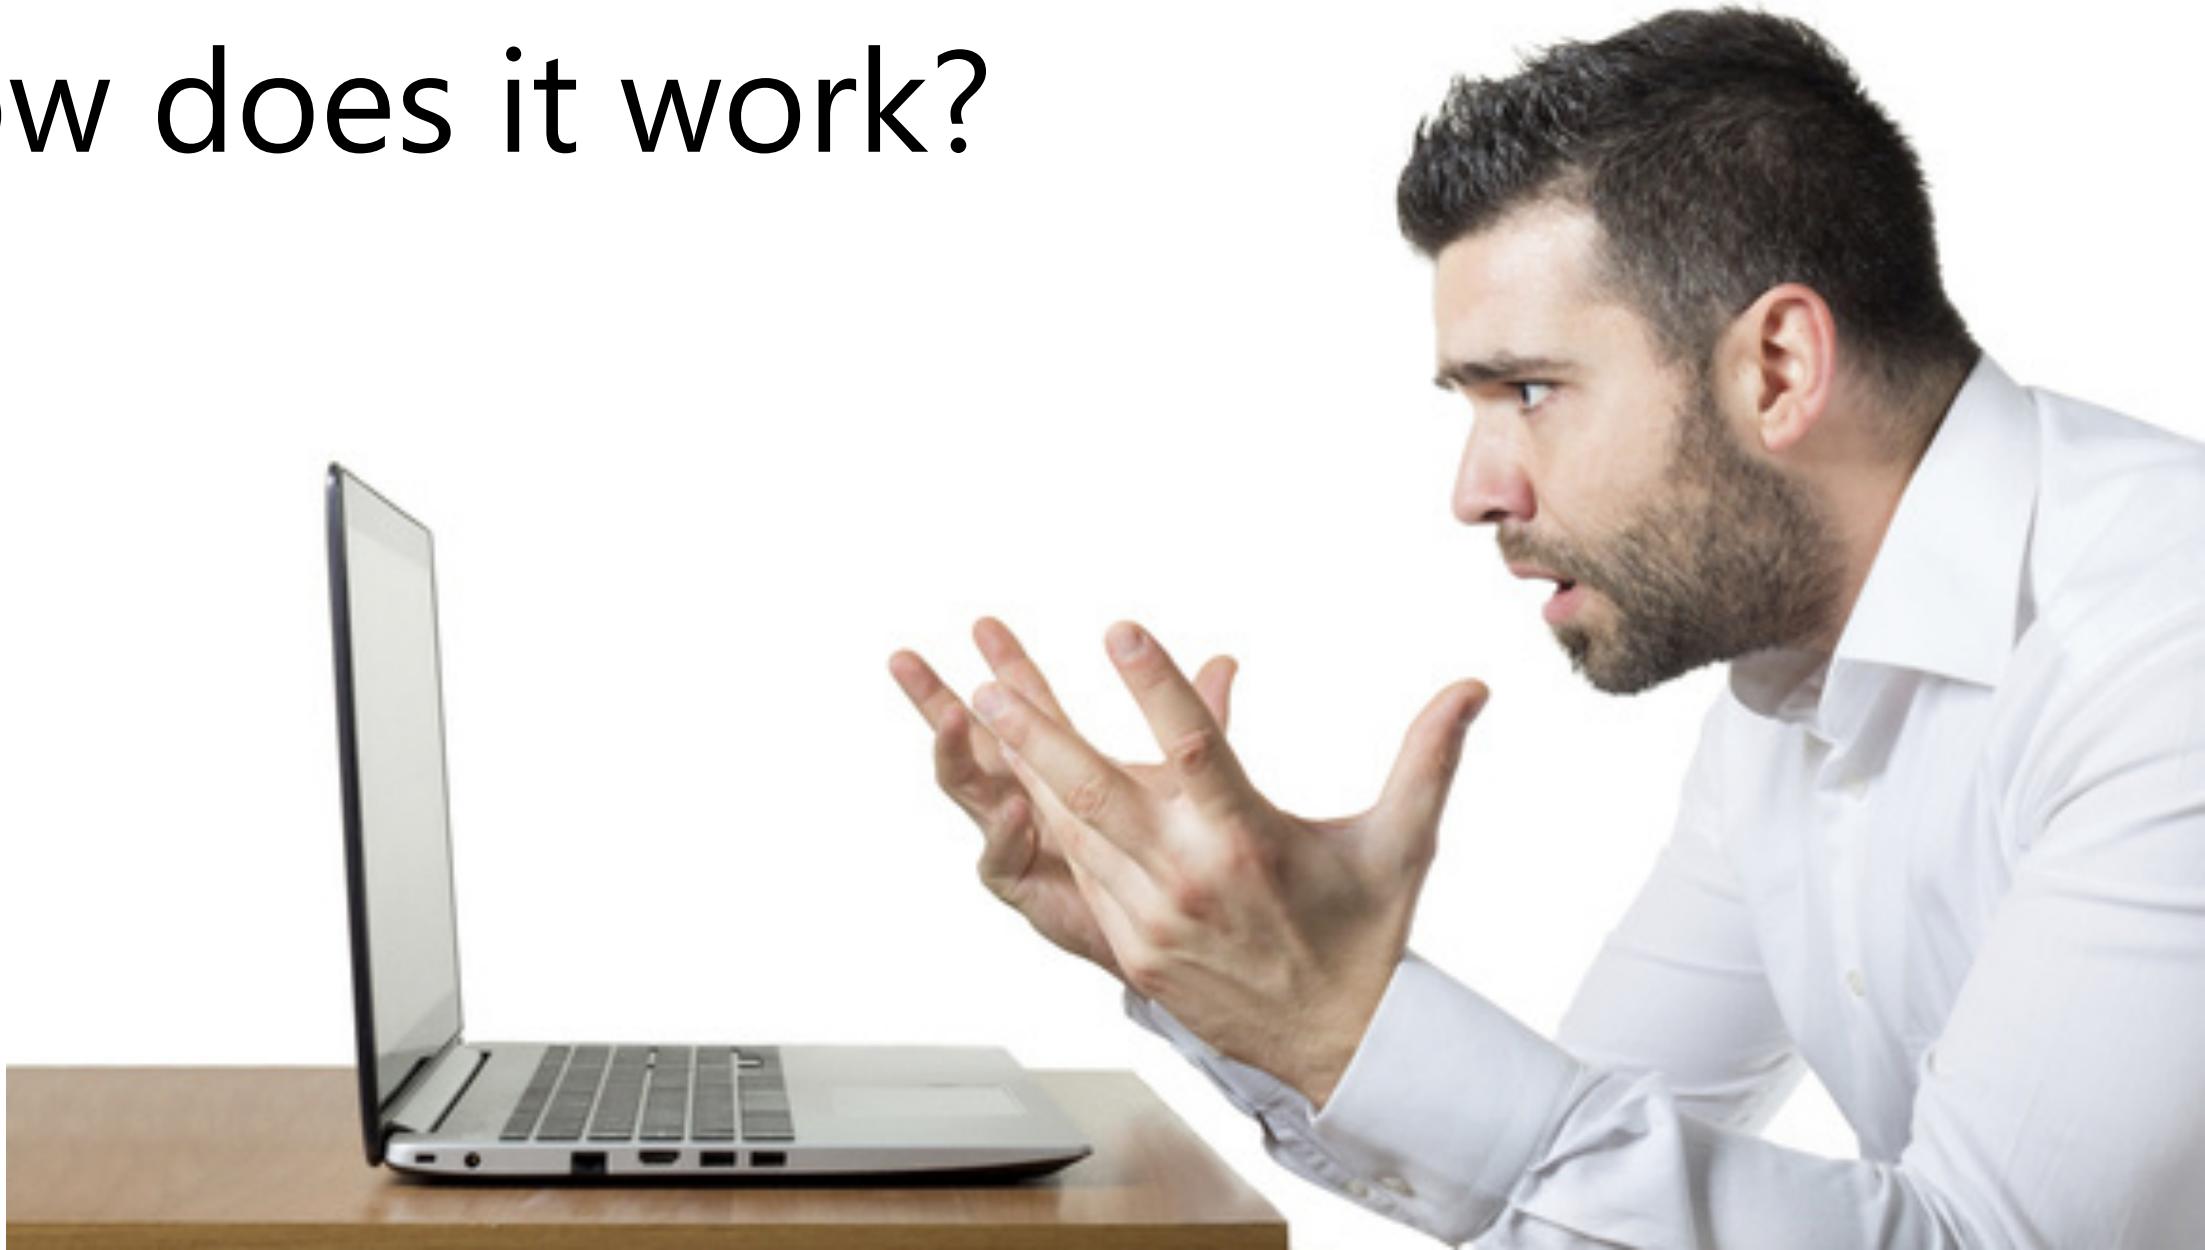


Understanding Service Mesh

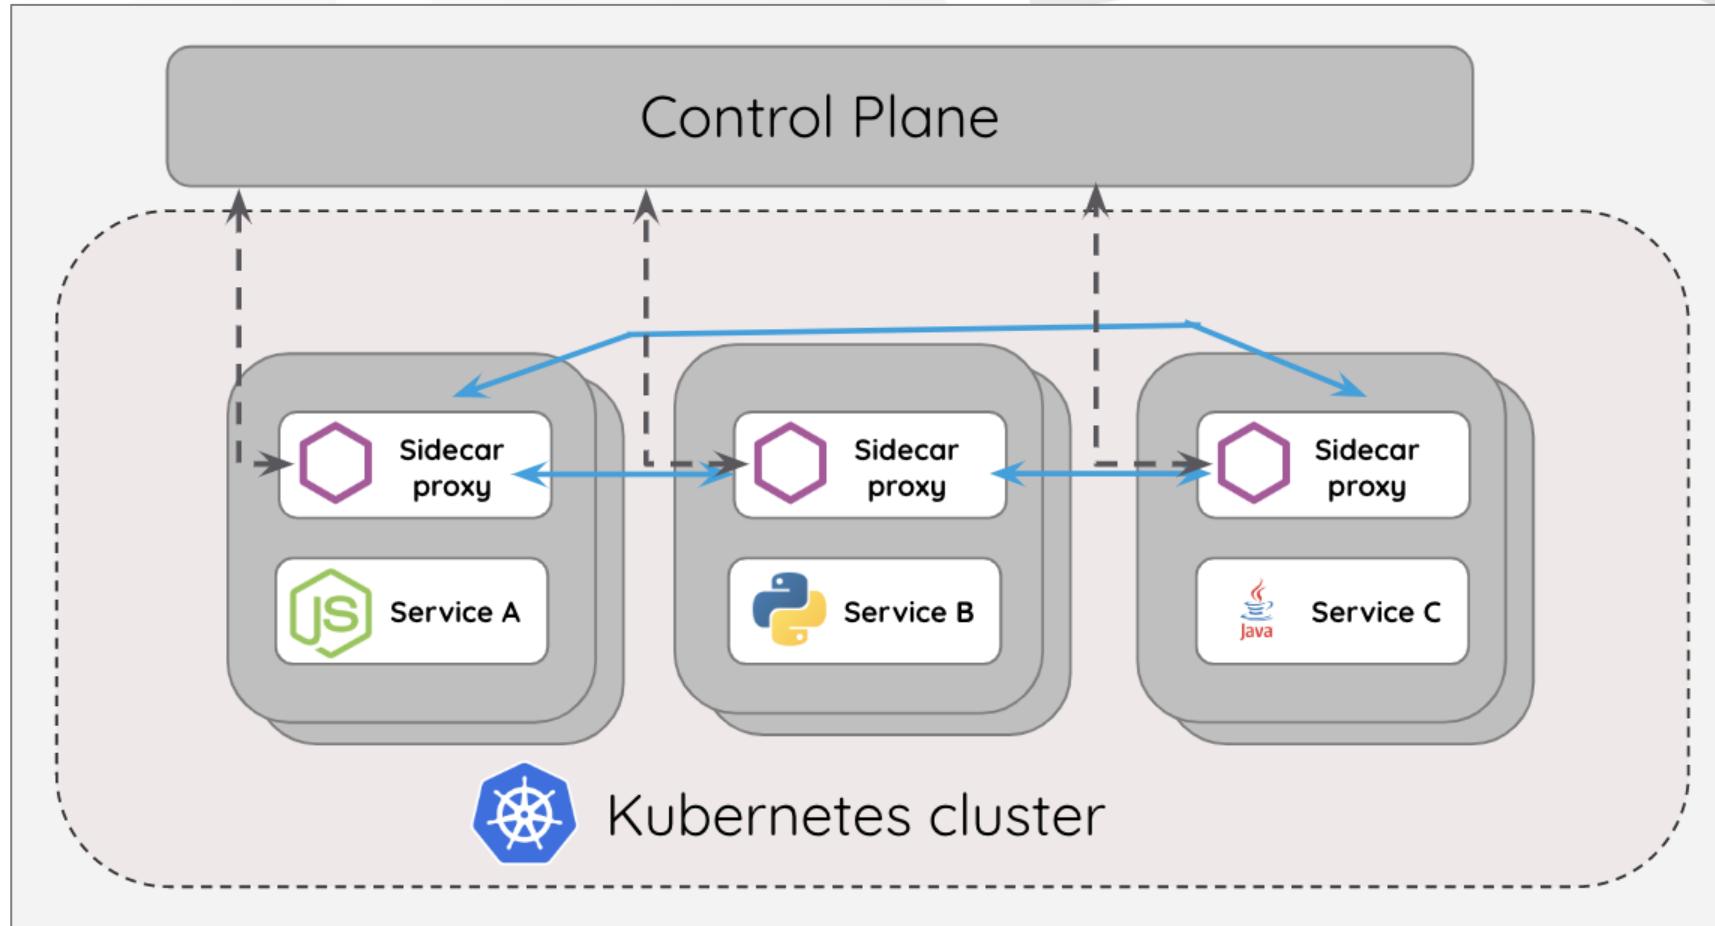
Observability → Visibility, Monitoring, Tracing, etc



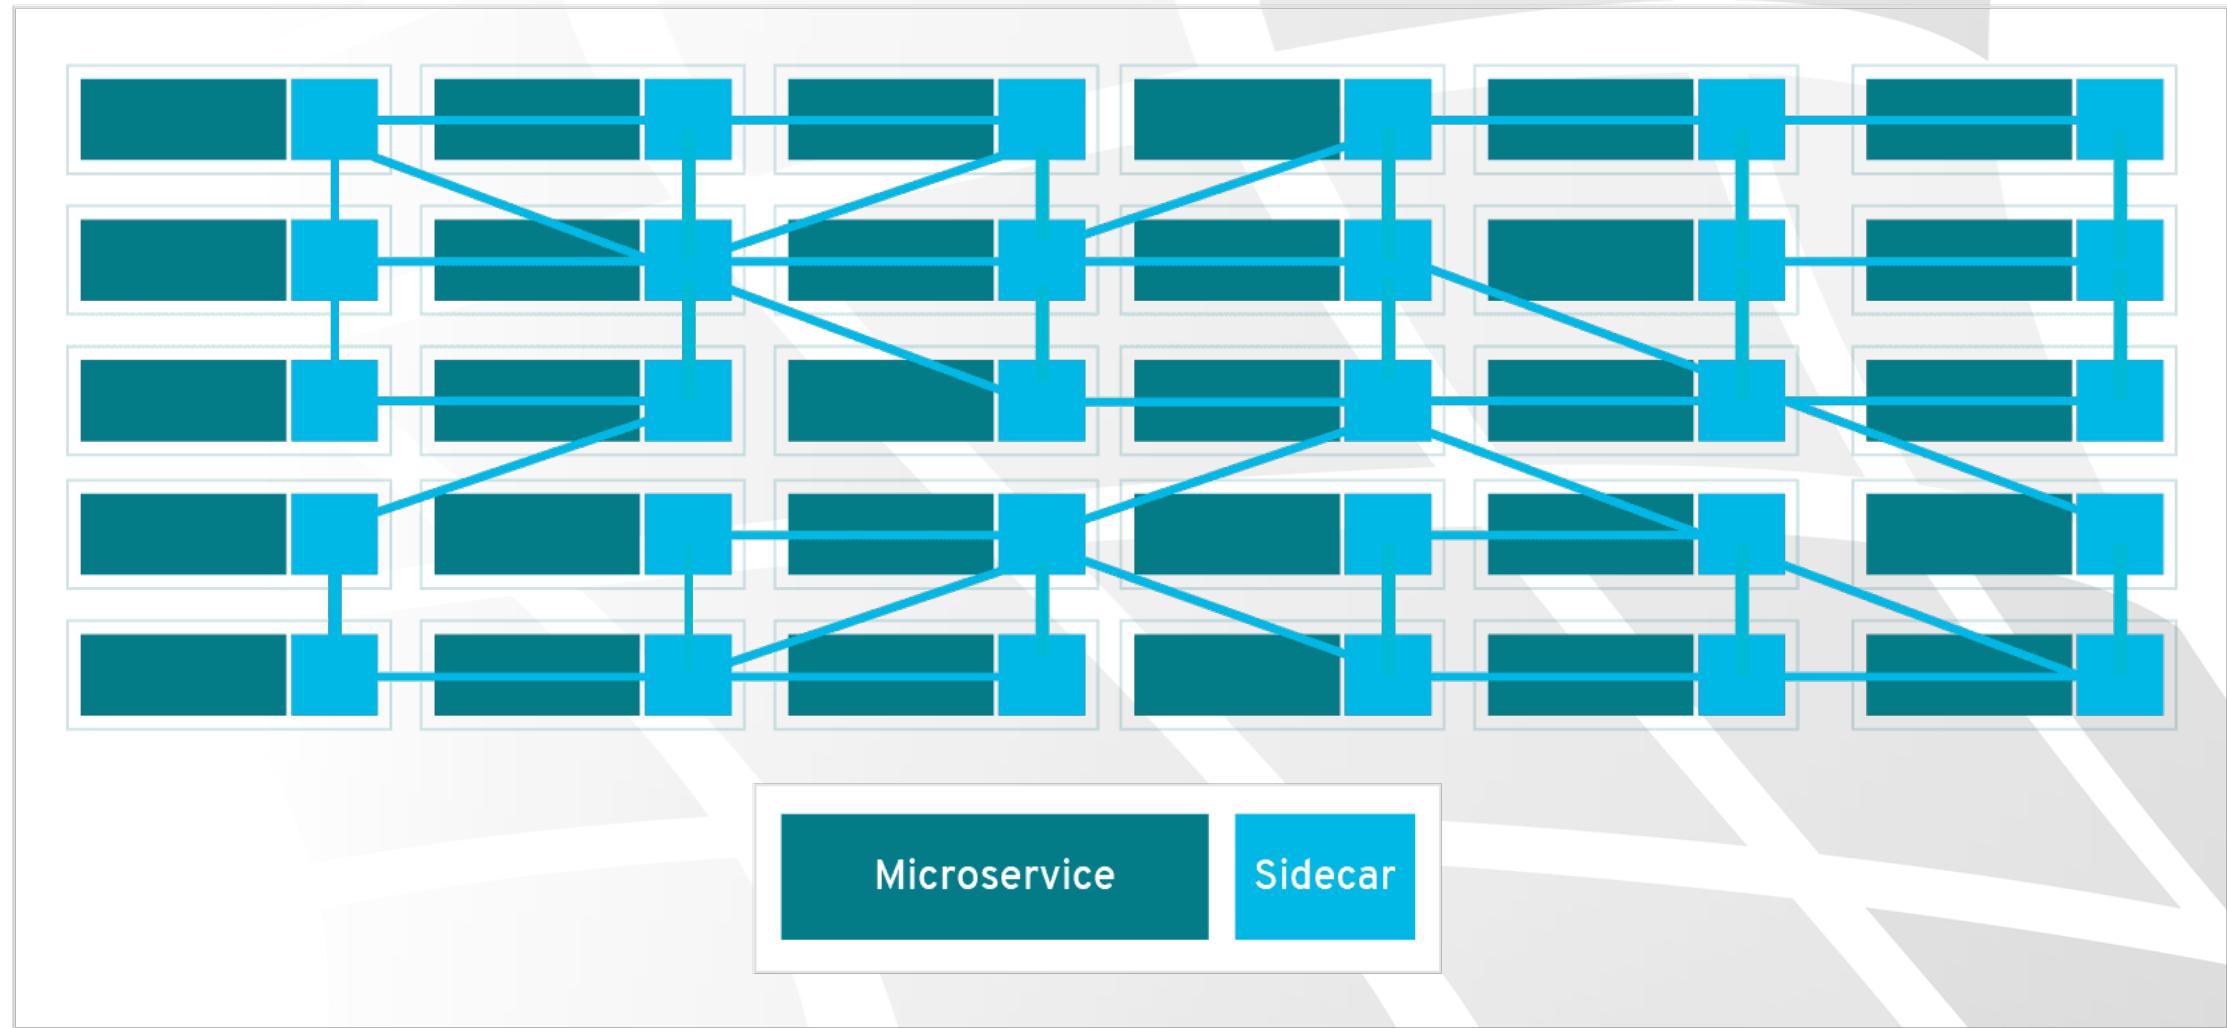
How does it work?



Understanding Service Mesh

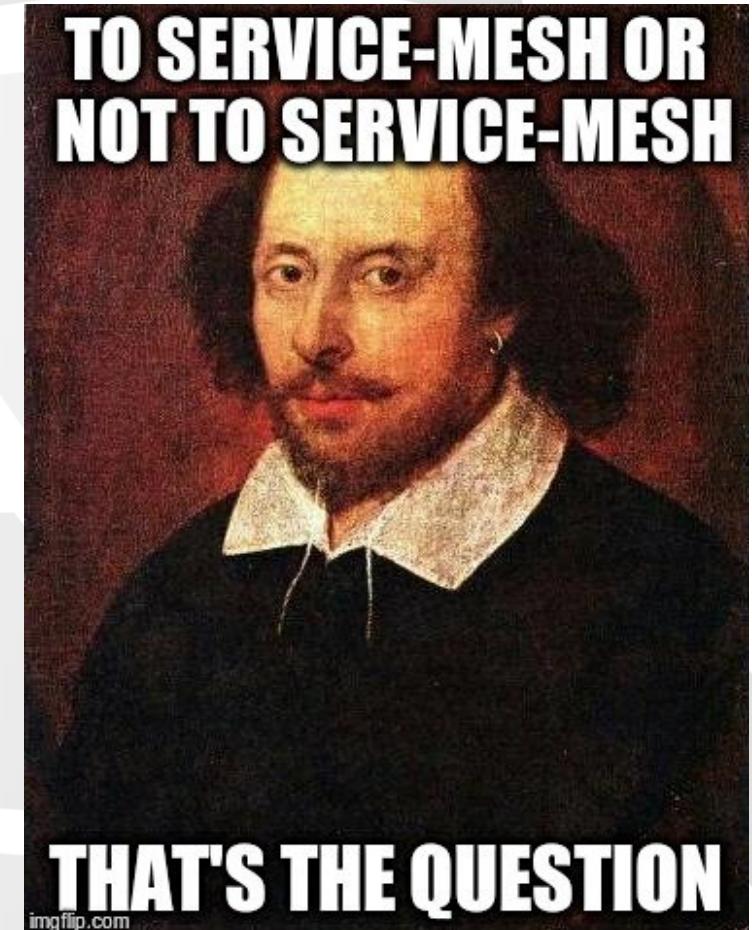


Understanding Service Mesh



Why use a Service Mesh?

- Automatic load balancing
- Fine-grained control of traffic
- Pluggable policy layer
- Access controls, rate limits and quotas.
- Service discovery
- Service monitoring with automatic metrics, logs and traces for all traffic.
- Secure service to service communication



Service Mesh Implementations



Istio

If you have heard about service mesh, you have probably heard about Istio too. Istio is by far the most popular service mesh because of its rich feature set and it Google's and IBM's support.



Linkerd 2

While Istio made the service mesh popular, Linkerd was the first service mesh and quite successful. Still, the developers decided to build a new version - Linkerd 2 - committed to usability, performance, and Kubernetes as the underlying platform.



Consul/Consul Connect

HashiCorp's Consul has been well known as a service discovery solution for a long time. Now that it has adopted the Envoy proxy and Sidecar pattern, Consul can serve as a service mesh for a variety of platforms like Kubernetes and VMs.



AWS App Mesh

Not long after the service mesh hype, AWS added its own service mesh for applications on AWS.



Traefik Mesh

As the name already reveals, Traefik Mesh (formerly Maesh) is the service mesh based on the cloud-native API gateway Traefik.



Kuma

Similar to Traefik Mesh, Kuma is also a very new service mesh made by developers of an API gateway - Kong.



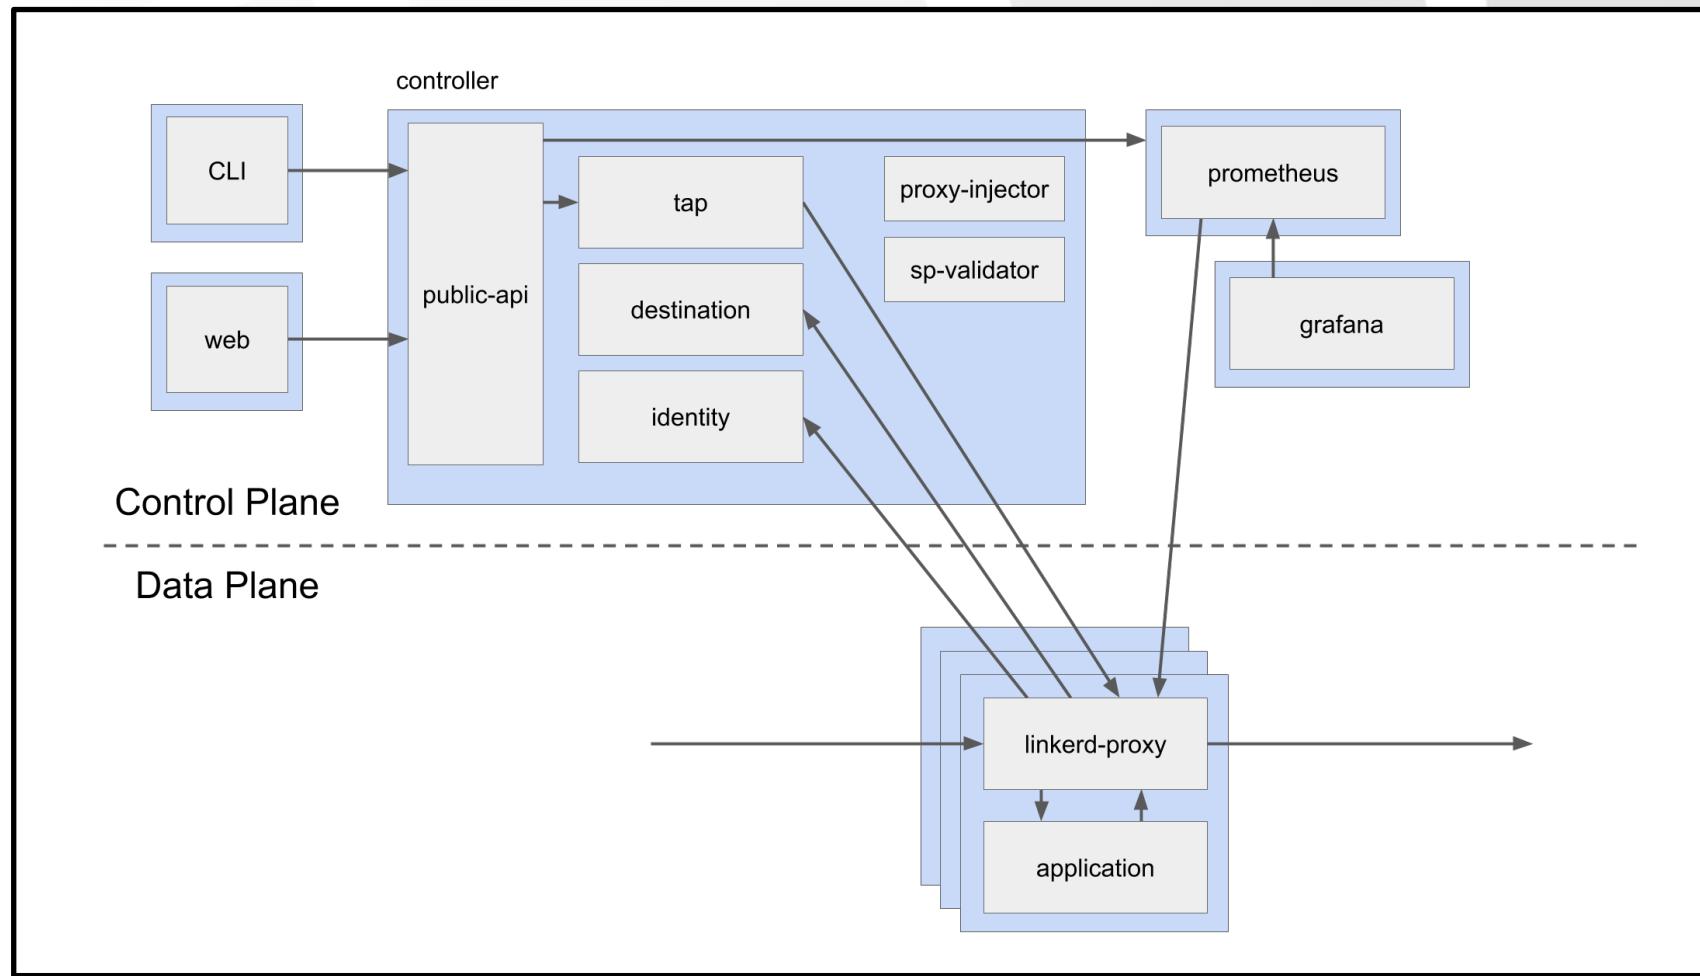
Open Service Mesh (OSM)

A new implementation by Microsoft, following common service mesh design principles like adopting envoy and implementing SMI spec.

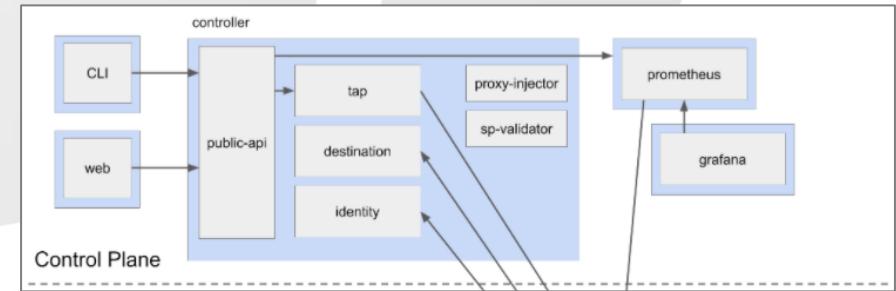
About Linkerd

- ★ Linkerd is focused on simplicity, speed, and low resource usage
- ★ Linkerd is built for security from the ground up (on-by-default mTLS, a data plane that is built in a memory-safe language, regular security audits, etc)
- ★ Linkerd is committed to open governance and is hosted by a neutral foundation (CNCF)
- ★ Use the `linkerd2-proxy`, which is built specifically for the service mesh sidecar use case, (smaller and faster than envoy-based service meshes)

Linkerd Architecture

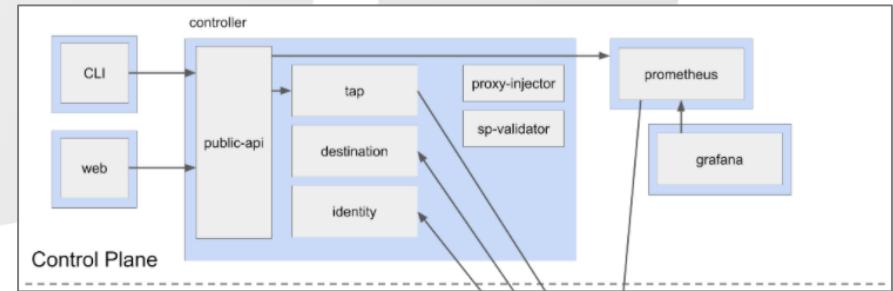


Architecture - Control Plane



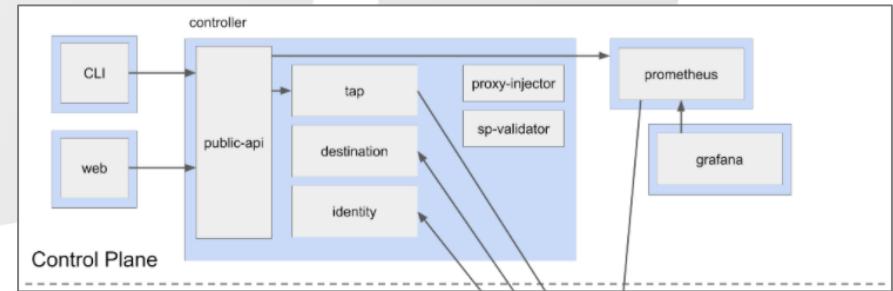
- ★ **Web:** The web deployment provides the Linkerd dashboard. This does not require running Linkerd dashboard and can be exposed to others.
- ★ **Public API:** The controller deployment consists of the public-api container that provides an API for the CLI and dashboard to interface with.
- ★ **Tap:** The tap deployment receives requests from the CLI and dashboard to watch requests and responses in real time.

Architecture - Control Plane



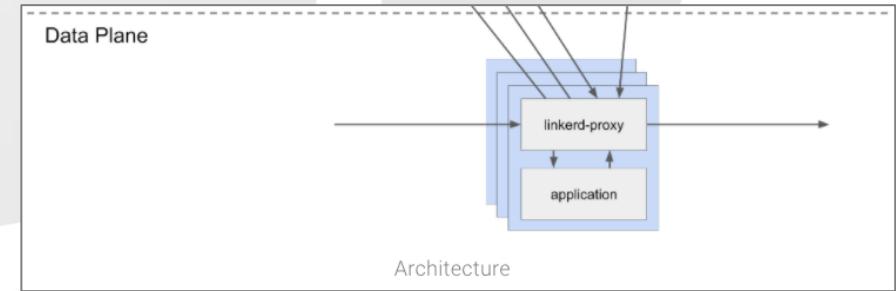
- ★ **Destination:** Each proxy in the data plane uses this component to lookup where to send requests.
- ★ **Identity:** Provides a Certificate Authority that accepts CSRs from proxies and returns certificates signed with the correct identity.
- ★ **Proxy Injector:** The injector is an admission controller, which receives a webhook request every time a pod is created.

Architecture - Control Plane



- ★ **Service Profile Validator:** The validator is also an admission controller, which validates new Service Profiles (Linkerd CRD) before they are saved.
- ★ **Grafana:** Linkerd comes with many dashboards out of the box. The Grafana component is used to render and display these dashboards.
- ★ **Prometheus:** Prometheus is a cloud native monitoring solution that is used to collect and store all of the Linkerd metrics.

Architecture - Data Plane



★ **Proxy:** An ultralight transparent proxy written in Rust, the proxy is installed into each pod of a service and becomes part of the data plane. It receives all incoming traffic for a pod and intercepts all outgoing traffic via an initContainer that configures iptables to forward the traffic correctly. Because it is a sidecar and intercepts all the incoming and outgoing traffic for a service, there are no code changes required and it can even be added to a running service.

Linkerd CLI

The Linkerd CLI is the primary way to interact with Linkerd. It can install the control plane to your cluster, add the proxy to your service and provide detailed metrics for how your service is performing.

<code>check</code>	Check the Linkerd installation for potential problems
<code>completion</code>	Output shell completion code for the specified shell (bash or zsh)
<code>dashboard</code>	Open the Linkerd dashboard in a web browser
<code>edges</code>	Display connections between resources, and Linkerd proxy identities
<code>endpoints</code>	Introspect Linkerd's service discovery state
<code>get</code>	Display one or many mesh resources
<code>inject</code>	Add the Linkerd proxy to a Kubernetes config
<code>install</code>	Output Kubernetes configs to install Linkerd
<code>install-cni</code>	Output Kubernetes configs to install Linkerd CNI
<code>install-sp</code>	Output Kubernetes configs to install Linkerd Service Profiles

<code>logs</code>	Tail logs from containers in the Linkerd control plane
<code>metrics</code>	Fetch metrics directly from Linkerd proxies
<code>profile</code>	Output service profile config for Kubernetes
<code>routes</code>	Display route stats
<code>stat</code>	Display traffic stats about one or many resources
<code>tap</code>	Listen to a traffic stream
<code>top</code>	Display sorted information about live traffic
<code>uninject</code>	Remove the Linkerd proxy from a Kubernetes config
<code>upgrade</code>	Output Kubernetes configs to upgrade an existing Linkerd control plane
<code>version</code>	Print the client and server version information

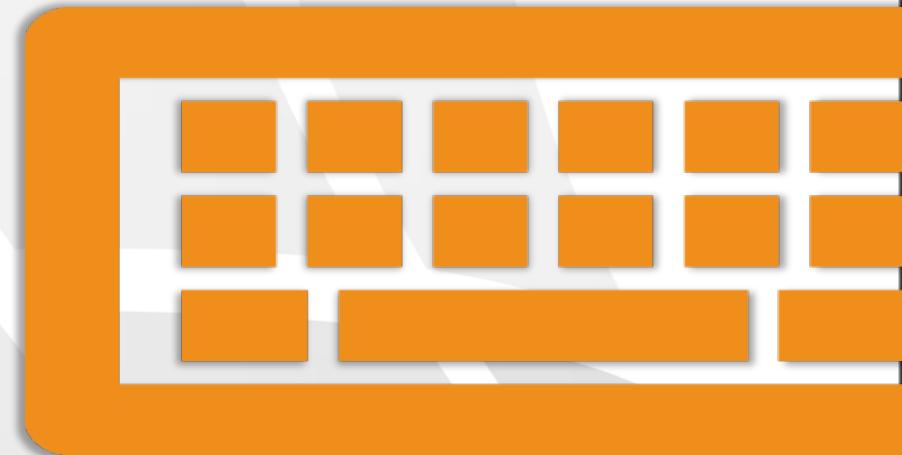
Linkerd Installation



- ★ Ensure you have access to a Kubernetes cluster running 1.13 or later
- ★ Ensure you have a functioning kubectl command on your local machine
- ★ Install the Linkerd CLI
- ★ Validate your Kubernetes Cluster
- ★ Install Linkerd onto the Cluster
- ★ Validate the Installation

Lab 01: Installing Linkerd

Lab



<https://github.com/leonjalfon1/sdp-servicemesh/labs/01-installing-linkerd.md>

Linkerd Dashboard

The Linkerd dashboard provides a high-level view of what is happening with your services in real time.

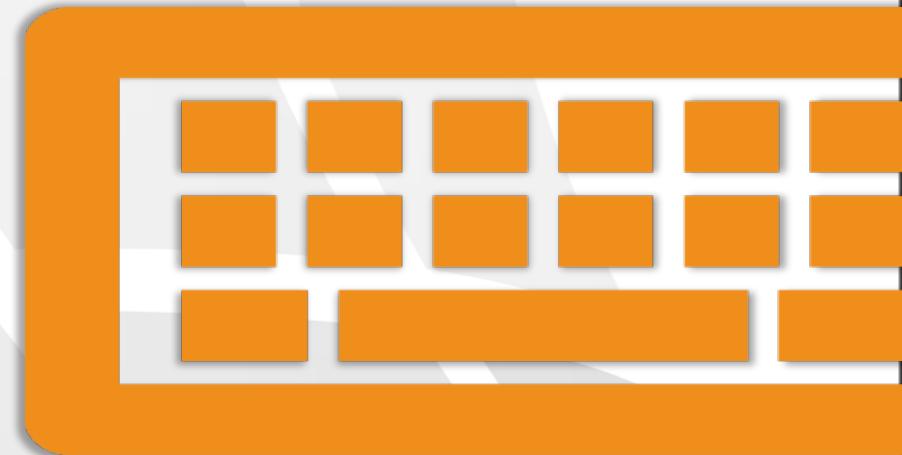
The screenshot shows the Linkerd dashboard interface. On the left is a sidebar with navigation links for CLUSTER (Namespaces, Control Plane) and WORKLOADS (Daemon Sets, Deployments, Jobs, Pods, Replication Controllers). A dropdown menu shows 'EMOJIVOTO' is selected. The main area has two sections: 'Deployments' and 'Pods'. The 'Deployments' section lists four services: emoji, vote-bot, voting, and web. The 'Pods' section lists two corresponding pods: emoji-697b575bd9-9kkcf and vote-bot-7bd97dfbdc-7psdm. Both sections include columns for Meshed, Success Rate, RPS, P50 Latency, P95 Latency, P99 Latency, and Grafana monitoring links.

Deployment	Meshed	Success Rate	RPS	P50 Latency	P95 Latency	P99 Latency	Grafana
emoji	1/1	100.00% ●	1.97	1 ms	1 ms	1 ms	
vote-bot	1/1	--	--	--	--	--	
voting	1/1	83.05% ●	0.98	1 ms	1 ms	1 ms	
web	1/1	89.83% ●	1.97	4 ms	9 ms	10 ms	

Pod	Meshed	Success Rate	RPS	P50 Latency	P95 Latency	P99 Latency	Grafana
emoji-697b575bd9-9kkcf	1/1	100.00% ●	1.97	1 ms	1 ms	1 ms	
vote-bot-7bd97dfbdc-7psdm	1/1	--	--	--	--	--	

Lab 02: Exploring Linkerd

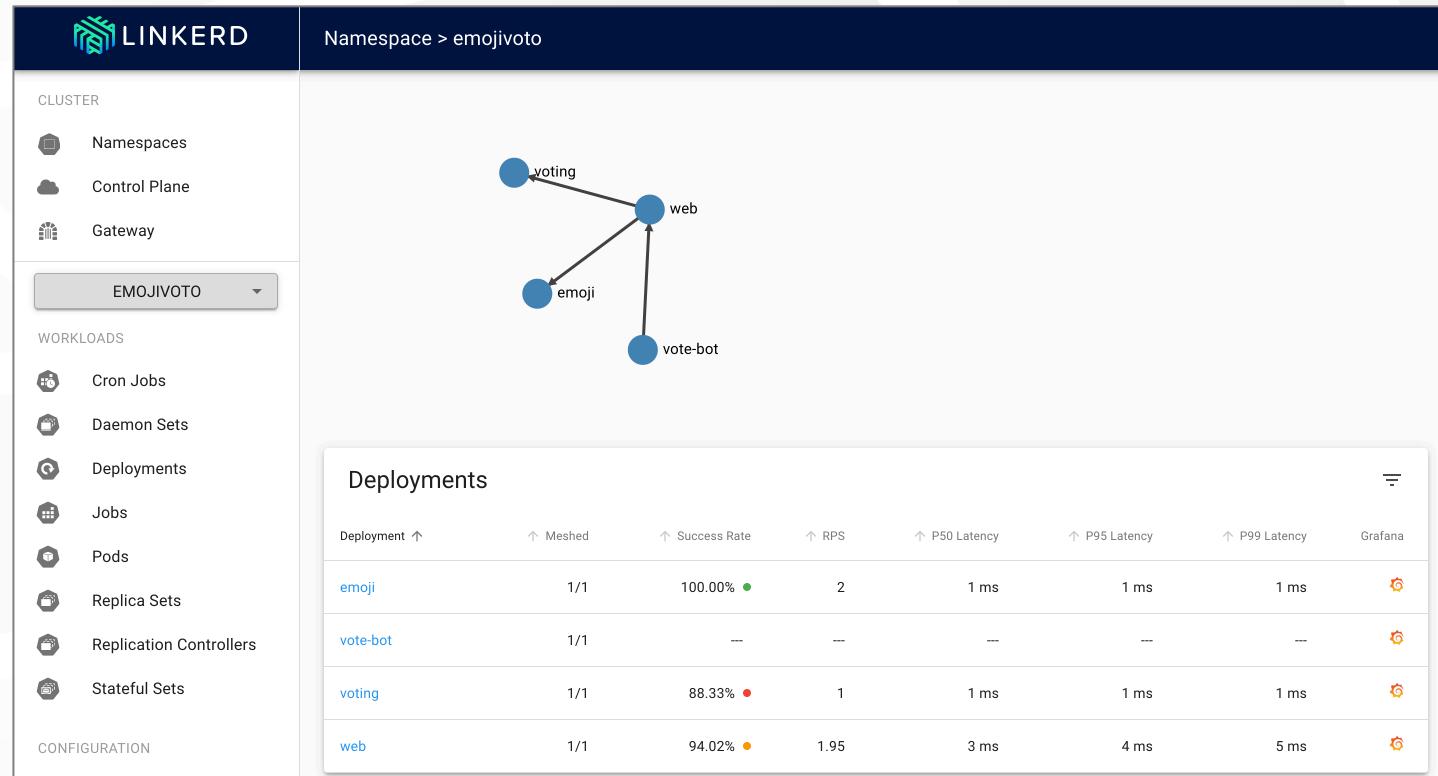
Lab



<https://github.com/leonjalfon1/sdp-servicemesh/labs/02-exploring-linkerd.md>

Linkerd Dashboard

The Linkerd dashboard provides a high-level view of what is happening with your services in real time.



Linkerd Dashboard

Show the “golden” metrics for each deployment

The screenshot shows the Linkerd Dashboard interface. On the left, there's a sidebar with navigation links for CLUSTER (Namespaces, Control Plane) and WORKLOADS (Daemon Sets, Deployments, Jobs, Pods, Replication Controllers). A dropdown menu shows 'EMOJIVOTO'. The main area has two sections: 'Deployments' and 'Pods', each with a table of metrics.

Deployments

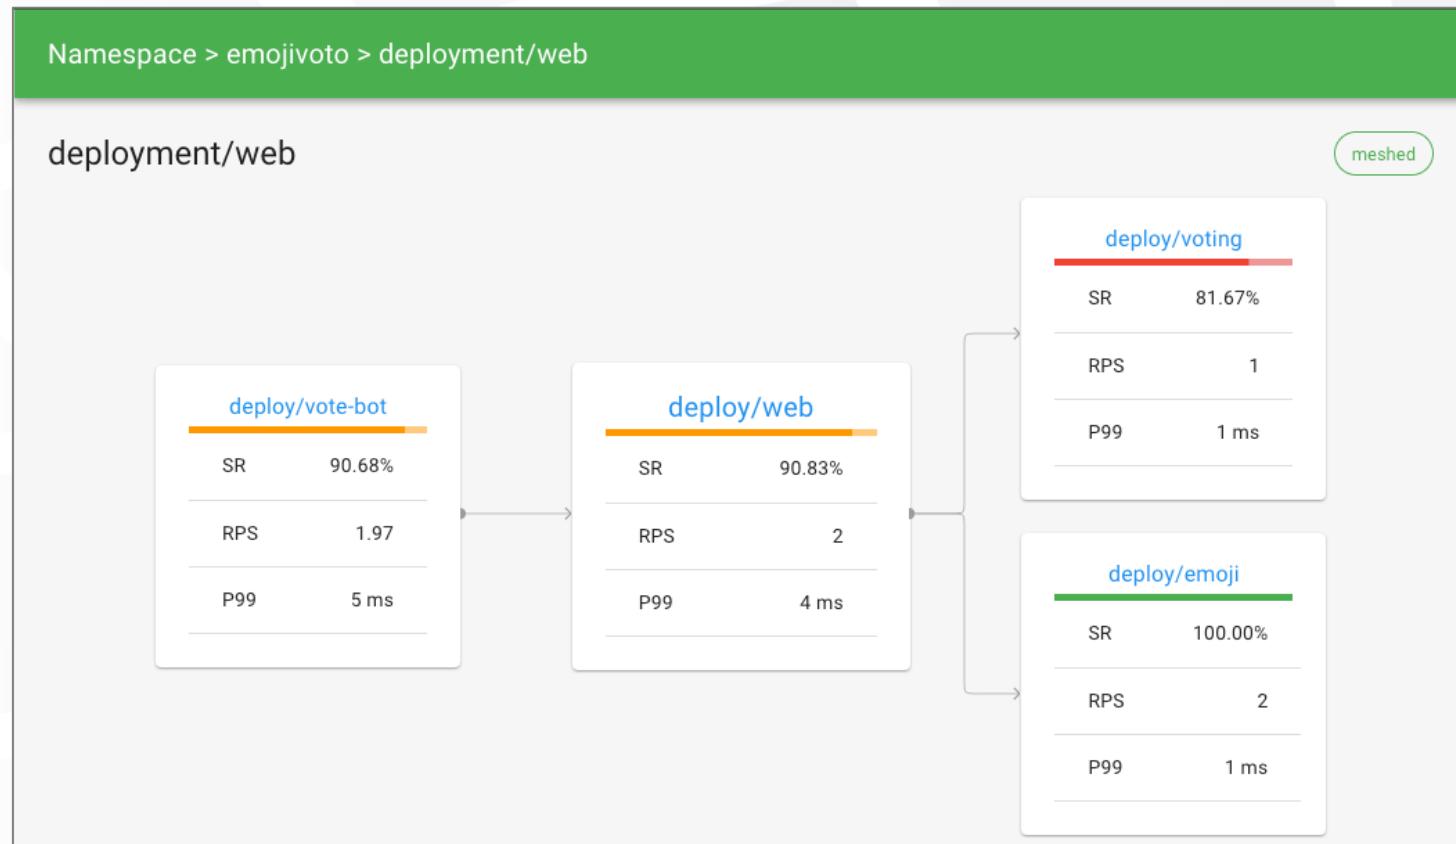
Deployment ↑	↑ Meshed	↑ Success Rate	↑ RPS	↑ P50 Latency	↑ P95 Latency	↑ P99 Latency	Grafana
emoji	1/1	100.00% ●	1.97	1 ms	1 ms	1 ms	
vote-bot	1/1	--	--	--	--	--	
voting	1/1	83.05% ●	0.98	1 ms	1 ms	1 ms	
web	1/1	89.83% ●	1.97	4 ms	9 ms	10 ms	

Pods

Pod ↑	↑ Meshed	↑ Success Rate	↑ RPS	↑ P50 Latency	↑ P95 Latency	↑ P99 Latency	Grafana
emoji-697b575bd9-9kkcf	1/1	100.00% ●	1.97	1 ms	1 ms	1 ms	
vote-bot-7bd97dfbdc-7psdm	1/1	--	--	--	--	--	

Linkerd Dashboard

Show deployment details (visualization)



Linkerd Dashboard

Top: Display sorted information about live traffic on a per-route basis.

The screenshot shows the Linkerd Top dashboard interface. At the top, there are dropdown menus for 'Namespace' (emojivoto) and 'Resource' (deployment/web), and two buttons: 'START' and 'RESET'. Below this, a 'Current Top query' section displays the command: 'linkerd top deployment/web --namespace emojivoto'. The main area is a table titled 'Top' with columns: Name, Method, Path, Count, Best, Worst, Last, Success Rate, and Tap. The table lists seven routes, each with a blue icon and a link icon:

	Name	Method	Path	Count	Best	Worst	Last	Success Rate	Tap
TO	deploy/emoji	POST	/emojivoto.v1.EmojiService/ListAll	9	1 ms	2 ms	1 ms	100.00% ●	🔗
FROM	deploy/vote-bot	GET	/api/list	9	3 ms	4 ms	3 ms	100.00% ●	🔗
TO	deploy/emoji	POST	/emojivoto.v1.EmojiService/FindByShortcode	9	787 µs	2 ms	940 µs	100.00% ●	🔗
FROM	deploy/vote-bot	GET	/api/vote	9	3 ms	5 ms	4 ms	88.89% ●	🔗
TO	deploy/voting	POST	/emojivoto.v1.VotingService/VoteBrideWithVeil	2	1 ms	1 ms	1 ms	100.00% ●	🔗
TO	deploy/voting	POST	/emojivoto.v1.VotingService/VoteWoman	1	1 ms	1 ms	1 ms	100.00% ●	🔗
TO	deploy/voting	POST	/emojivoto.v1.VotingService/VoteWoman	1	0 ms	0 ms	0 ms	100.00% ●	🔗

Linkerd Dashboard

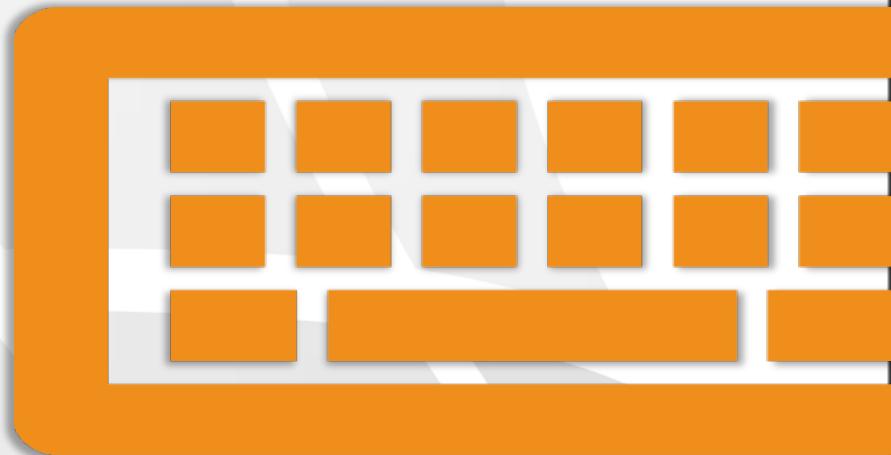
Tap: Listen to a traffic stream for a resource

The screenshot shows the Linkerd Tap interface. At the top, there are dropdown menus for 'Namespace' (emojivoto) and 'Resource' (deployment/web), along with 'START' and 'RESET' buttons. Below these, a 'Current Tap query' field contains the command: `linkerd tap deployment/web --namespace emojivoto`. A 'Show more filters' button is also present. The main area displays a table of network traffic logs:

Direction	Name	Method	Path	Latency	HTTP status	GRPC status
FROM	deploy/vote-bot	GET	/api/list	3 ms	200	--
FROM	deploy/vote-bot	GET	/api/vote	5 ms	500	--
TO	deploy/emoji	POST	/emojivoto.v1.EmojiService/FindByShortcode	2 ms	200	OK
TO	deploy/voting	POST	/emojivoto.v1.VotingService/VoteDoughnut	1 ms	200	Unknown
TO	deploy/emoji	POST	/emojivoto.v1.EmojiService/ListAll	1 ms	200	OK
FROM	deploy/vote-bot	GET	/api/vote	5 ms	200	--

Lab 03: Debugging with Linkerd

Lab



<https://github.com/leonjalfon1/sdp-servicemesh/labs/03-debugging-with-linkerd.md>

Automatic Mutual TLS

- ★ By default, Linkerd automatically enables mutual Transport Layer Security (mTLS) for most TCP traffic between meshed pods, by establishing and authenticating secure, private TLS connections between Linkerd proxies.
- ★ Because the Linkerd control plane also runs on the data plane, this means that communication between Linkerd's control plane components are also automatically secured via mTLS
- ★ Not all traffic can be automatically mTLS'd, but it's easy to verify which traffic is (using `linkerd tap`)

Automatic Mutual TLS – Linkerd Edges

The screenshot shows the Linkerd dashboard interface. On the left, there's a sidebar with navigation links for CLUSTER (Namespaces, Control Plane, Gateway), LINKERD (selected), WORKLOADS (Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets), and CONFIGURATION.

The main content area has a header "Namespace > linkerd > deployment/linkerd-web". It displays three sections: Pod metrics, TCP metrics, and Edge configurations.

Pod Metrics:

Pod	Meshed	Success Rate	RPS	P50 Latency	P95 Latency	P99 Latency	Grafana
linkerd-web-56fdbfd448-hszpr	1/1	100.00% ●	1.43	211 ms	291 ms	298 ms	

TCP Metrics:

Pod	Meshed	Connections	Read Bytes / sec	Write Bytes / sec	Grafana
linkerd-web-56fdbfd448-hszpr	1/1	17	9.642kB/s	9.628kB/s	

Edges (Identity: linkerd-web.linkerd):

TO	Name	Identity	Secured
linkerd	linkerd-controller	linkerd-controller.linkerd	
linkerd	linkerd-grafana	linkerd-grafana.linkerd	
linkerd	linkerd-prometheus	linkerd-prometheus.linkerd	
linkerd	linkerd-tap	linkerd-prometheus.linkerd	

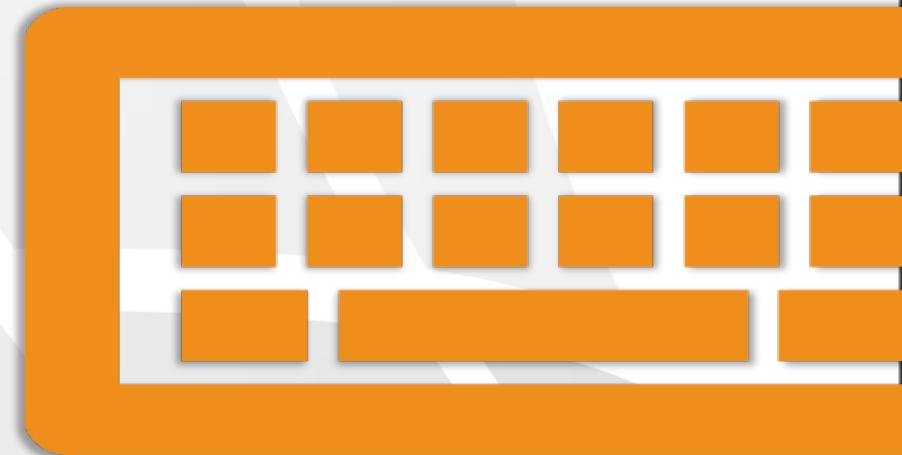
A red box highlights the "Secured" column, which contains green checkmarks for all four entries, indicating that automatic Mutual TLS is enabled for all edges.

Automatic Mutual TLS - Considerations

- ★ By default, the issuer certificate and key are not automatically rotated (expires after 365). You can set up automatic rotation with cert-manager
- ★ Linkerd does not currently enforce mTLS
- ★ Any unencrypted requests inside the mesh will be opportunistically upgraded to mTLS.
- ★ Any requests originating from inside or outside the mesh will not be automatically mTLS'd by Linkerd (this will be addressed in a future Linkerd release)

Lab 04: Validating Linkerd mTLS

Lab



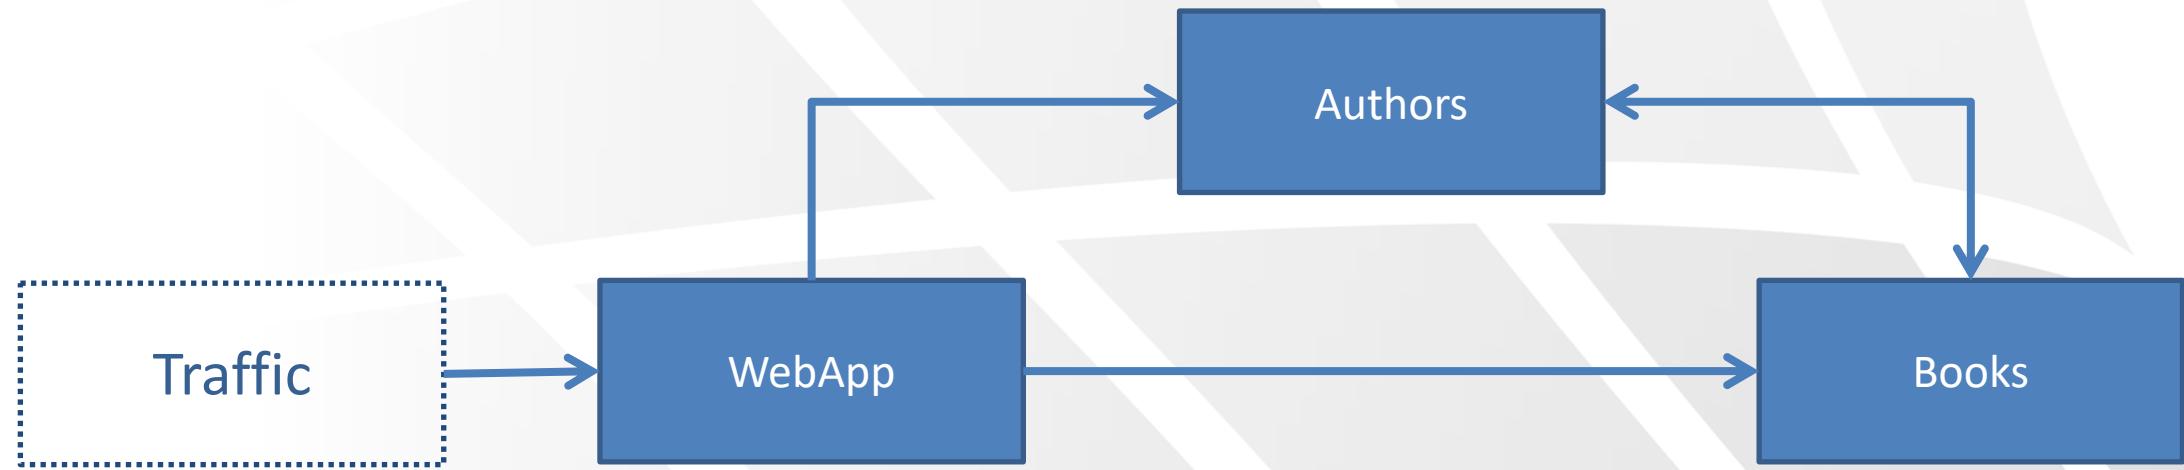
<https://github.com/leonjalfon1/sdp-servicemesh/labs/04-validating-linkerd-mtls.md>

BookApp Demo App

- ★ Let's use the BookApp to learn the following topics:
 - ★ Debug a network issue (a classic case of non-obvious, intermittent failure)
 - ★ Configure Linkerd Service Profiles
 - ★ Adding Retries
 - ★ Adding Timeouts

BookApp - Architecture

The BookApp is a Ruby application that helps you manage your bookshelf. It consists of multiple microservices and uses JSON over HTTP to communicate with the other services.



Note: for demo purposes, the app comes with a simple traffic generator

BookApp - Architecture

There is a bug in the app: if you click Add Book, it will fail 50% of the time

Buoyant Books App

Books (15)				
Title	Author	Pages		
1Q84	Haruki Murakami	1184	<input checked="" type="checkbox"/>	<input type="checkbox"/>
American Gods	Neil Gaiman	588	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Children of Men	P.D. James	241	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Cloud Atlas	David Mitchell	529	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Dear Life	Alice Munro	319	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Dispossessed, The	Ursula K. Le Guin	401	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Fifth Season, The	N.K. Jemisin	512	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Fire Next Time, The	James Baldwin	141	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Giovanni's Room	James Baldwin	176	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Kafka on the Shore	Haruki Murakami	436	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Left Hand of Darkness, The	Ursula K. Le Guin	304	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Mysteries of Pittsburgh, The	Michael Chabon	306	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Obelisk Gate, The	N.K. Jemisin	448	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Stone Sky, The	N.K. Jemisin	464	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Telegraph Avenue	Michael Chabon	468	<input checked="" type="checkbox"/>	<input type="checkbox"/>

Add a Book

Title

Author

Page count

Authors (11)				
Last Name	First Name	Book Count		
Atwood	Margaret	0	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Baldwin	James	2	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Chabon	Michael	2	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Cline	Ernest	0	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Gaiman	Neil	1	<input checked="" type="checkbox"/>	<input type="checkbox"/>
James	P.D.	1	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Jemisin	N.K.	3	<input checked="" type="checkbox"/>	<input type="checkbox"/>
K. Le Guin	Ursula	2	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Mitchell	David	1	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Munro	Alice	1	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Murakami	Haruki	2	<input checked="" type="checkbox"/>	<input type="checkbox"/>

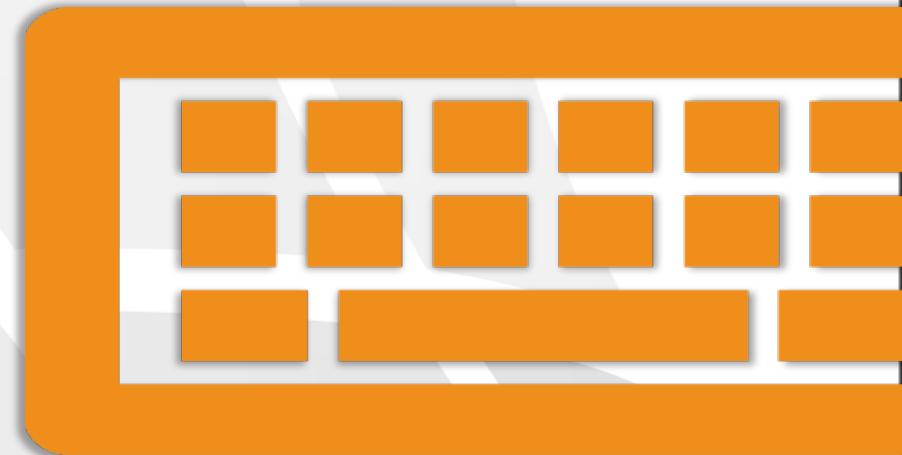
Add an Author

First name

Last name

Lab 05: Deploying and Debugging the BookApp

Lab



<https://github.com/leonjalfon1/sdp-servicemesh/labs/05-deploying-linkerd-bookapp.md>

Services Profiles

- ★ Is a custom Kubernetes resource (CRD) that can provide Linkerd additional information about a service.
- ★ Service Profiles allows you to define a list of routes for the service
- ★ Each route uses a regular expression to define which paths should match that route
- ★ Defining a service profile enables Linkerd to report per-route metrics and also allows you to enable per-route features such as retries and timeouts

Services Profiles

```
apiVersion: linkerd.io/v1alpha2
kind: ServiceProfile
metadata:
  creationTimestamp: null
  name: webapp.booksapp.svc.cluster.local
  namespace: booksapp
spec:
  routes:
    - condition:
        method: GET
        pathRegex: /
        name: GET /
    - condition:
        method: POST
        pathRegex: /authors
        name: POST /authors
```

Services Profiles

- ★ There are a couple different ways to use the linkerd profile command to create service profiles:
 - ★ **Swagger:** from an OpenAPI (Swagger) spec for your service
 - ★ **Protobuf:** from a protobuf format for your service
 - ★ **Auto-Creation:** generate service profiles from watching live traffic (this is based off tap data)
 - ★ **Template:** you can get a template that allows you to add routes manually

Services Profiles

★ Before using a Service Profile:

```
linkerd -n booksapp tap deploy/webapp -o wide | grep req
```

```
req id=124:12 proxy=in src=10.36.1.28:52398 dst=10.36.2.18:7000 tls=true :method=POST  
:authority=webapp:7000 :path=/books src_res=deploy/traffic src_ns=booksapp dst_res=deploy/webapp  
dst_ns=booksapp
```

★ After using a Service Profile:

```
linkerd -n booksapp tap deploy/webapp -o wide | grep req
```

```
req id=0:1 proxy=in src=10.1.3.76:57152 dst=10.1.3.74:7000 tls=true :method=POST  
:authority=webapp.default:7000 :path=/books/2878/edit src_res=deploy/traffic src_ns=booksapp  
dst_res=deploy/webapp dst_ns=booksapp rt_route=POST /books/{id}/edit
```

Services Profiles - Retries

- ★ To add a retry just add “isRetryable” to a specific route

```
spec:  
  routes:  
    - condition:  
        method: HEAD  
        pathRegex: /authors/[^/]*.json  
        name: HEAD /authors/{id}.json  
        isRetryable: true ### ADD THIS LINE ###
```

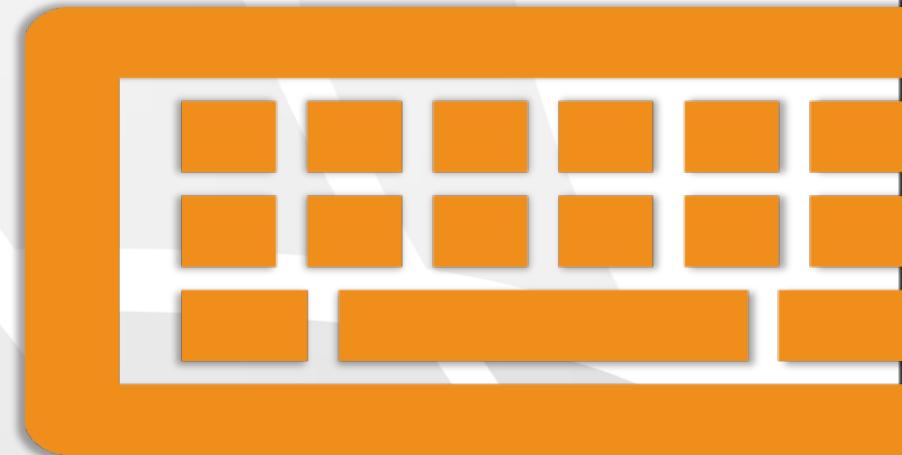
Services Profiles - Timeouts

- ★ To add a retry just add “timeout” to a specific route

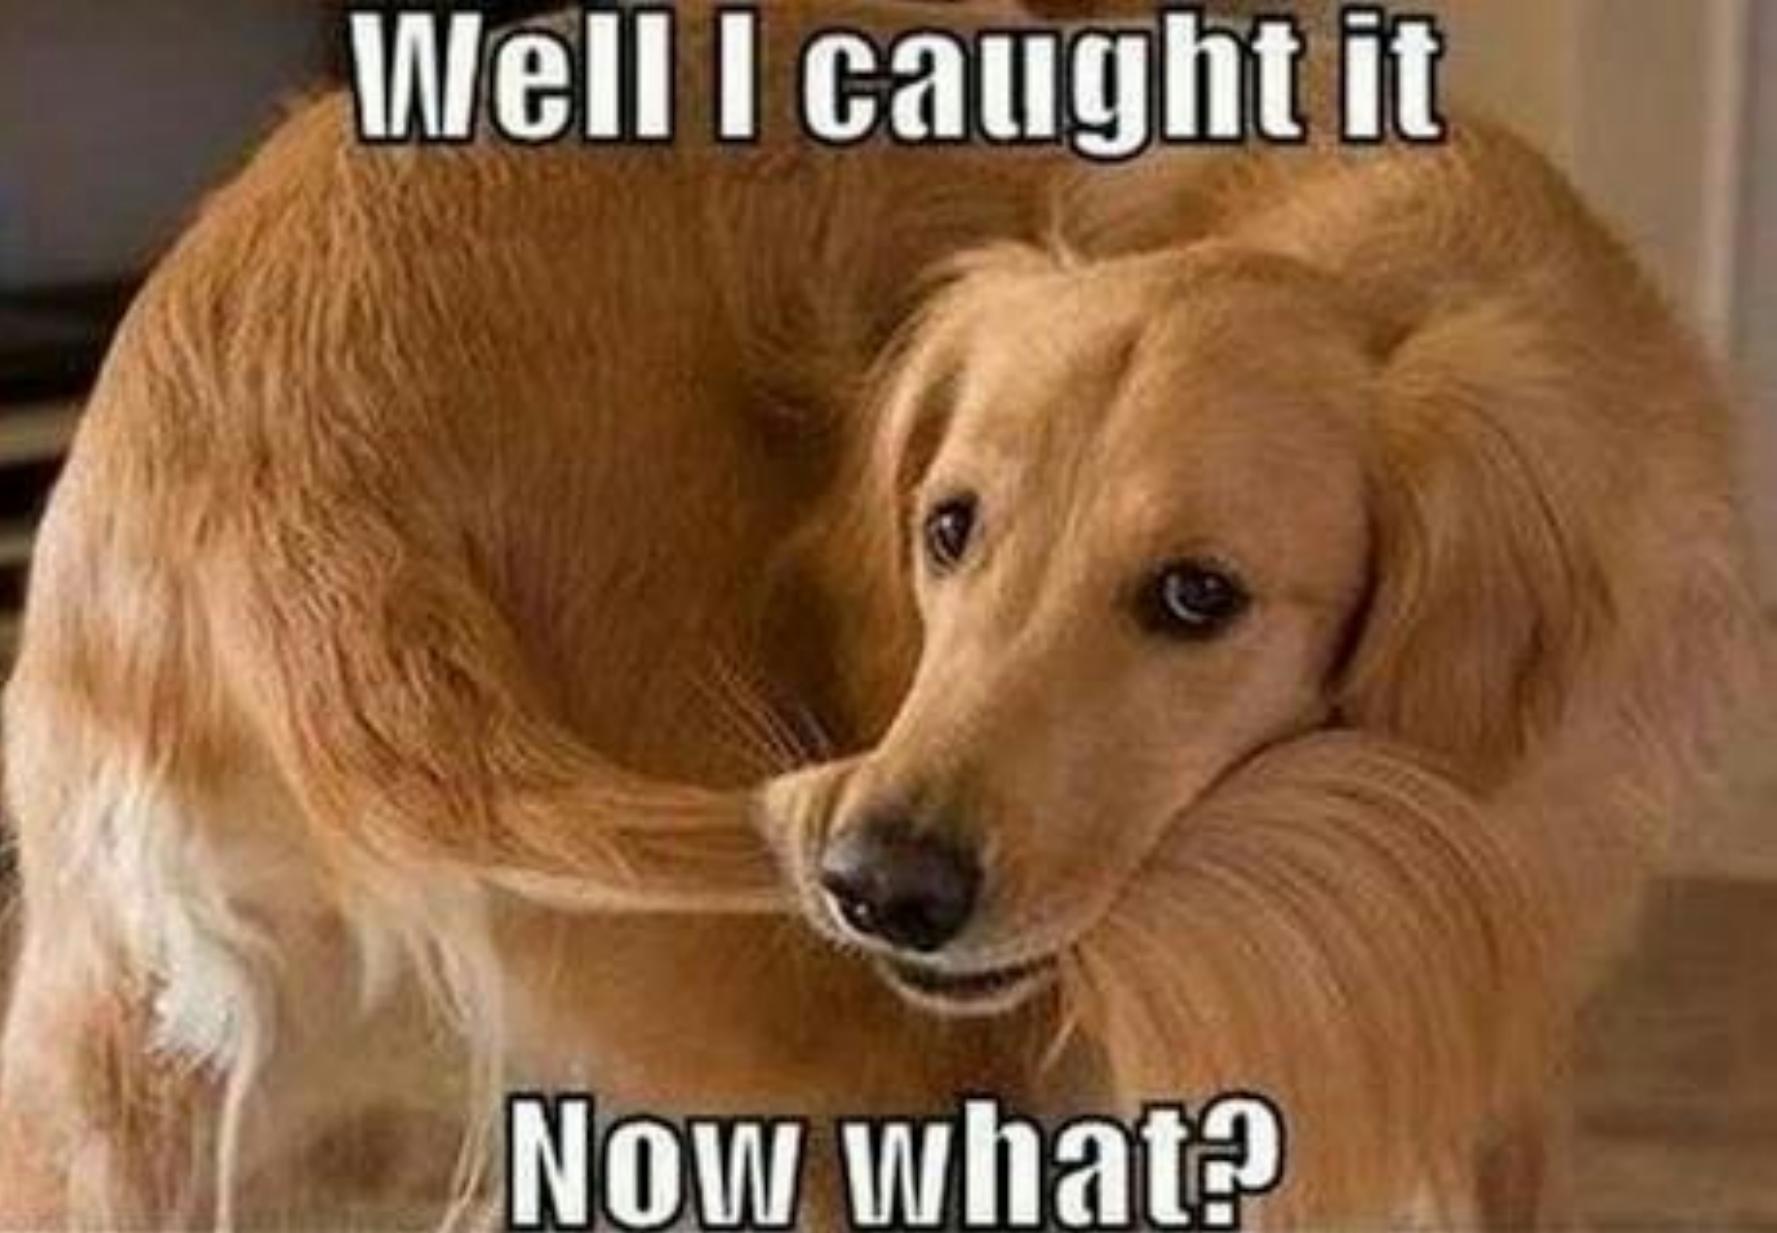
```
spec:  
  routes:  
    - condition:  
        method: PUT  
        pathRegex: /books/[^/]*\.json  
        name: PUT /books/{id}.json  
        timeout: 25ms ### ADD THIS LINE ###
```

Lab 06: Working with Service Profiles

Lab



<https://github.com/leonjalfon1/sdp-servicemesh/labs/06-linkerd-service-profiles.md>



Well I caught it

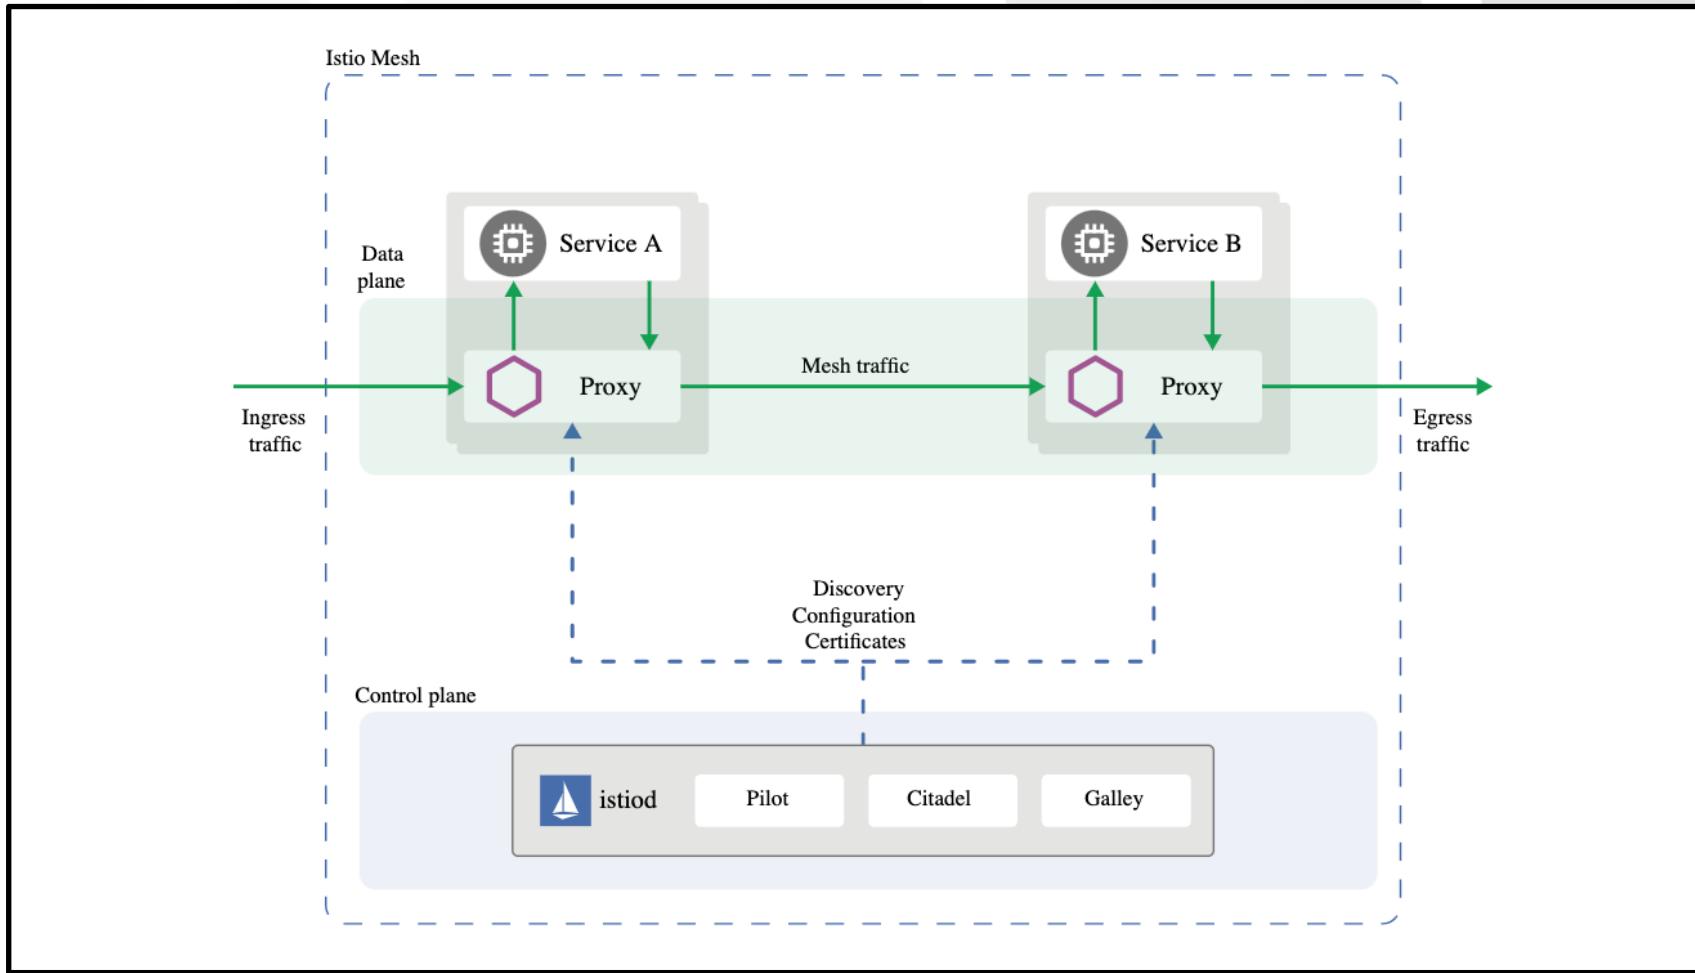
Now what?

About Istio

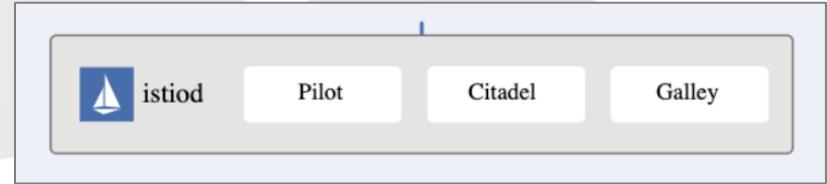


- ★ Initiative from Google, IBM and Lyft
- ★ Istio is designed to be platform-independent, initially focused on K8S.
- ★ Provides a uniform way to connect, manage, and secure microservices
- ★ It supports managing traffic flows between services, enforcing access policies, and aggregating telemetry data, all without requiring changes to the microservice code
- ★ Istio is totally open source and free

Istio Architecture

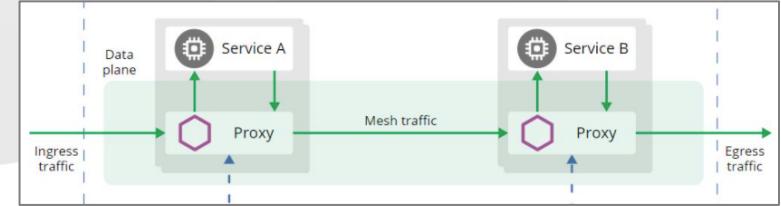


Architecture - Control Plane



- ★ **Istiod:** Unifies functionality that Pilot, Galley, Citadel and the sidecar injector previously performed, into a single binary (introduced in Istio v1.5)
- ★ **Pilot:** Provides service discovery, intelligent routing and fault tolerance
- ★ **Citadel:** Provides reliable service-to-service and end-user authentication.
- ★ **Galley:** Is responsible for validate the istio configuration.

Architecture - Data Plane



- ★ The data plane is composed of a set of intelligent proxies (Envoy) deployed as sidecars. These proxies mediate and control all network communication between microservices. They also collect and report telemetry on all mesh traffic.
- ★ Envoy is a CNCF Graduate based L4/L7 proxy developed on C++ by Lyft. Use API Driven updates (without hot-reload) and is injected as sidecar.

Istioctl

Istio configuration command line utility for service operators to debug and diagnose their Istio mesh.

```
Istio configuration command line utility for service operators to
debug and diagnose their Istio mesh.

Usage:
  istioctl [command]

Available Commands:
  analyze      Analyze Istio configuration and print validation messages
  authz        (authz is experimental. Use `istioctl experimental authz`)
  bug-report   Cluster information and log capture support tool.
  convert-ingress Convert Ingress configuration into Istio VirtualService configuration [Deprecated, it will be removed in Istio 1.9]
  dashboard    Access to Istio web UIs
  deregister   De-registers a service instance [Deprecated, it will be removed in Istio 1.9]
  experimental Experimental commands that may be modified or deprecated
  help         Help about any command
  install      Applies an Istio manifest, installing or reconfiguring Istio on a cluster.
  kube-inject  Inject Envoy sidecar into Kubernetes pod resources
  manifest     Commands related to Istio manifests
  operator     Commands related to Istio operator controller.
  profile      Commands related to Istio configuration profiles
  proxy-config Retrieve information about proxy configuration from Envoy [kube only]
  proxy-status Retrieves the synchronization status of each Envoy in the mesh [kube only]
  register     Registers a service instance (e.g. VM) joining the mesh [Deprecated, it will be removed in Istio 1.9]
  upgrade      Upgrade Istio control plane in-place
  validate     Validate Istio policy and rules files
  verify-install Verifies Istio Installation Status
  version      Prints out build version information
```

Istio Installation



- ★ **istioctl install:** The simplest and most qualified installation and management path with high security.
- ★ **Istio Operator:** Simple installation path without istioctl binaries.
- ★ **Istioctl Manifest Generate:** This method is suitable where strict auditing or augmentation of output manifests is needed.
- ★ **Helm installation:** The Helm charts used in this guide are the same underlying charts used when installing Istio via Istioctl or the Operator.

Istio Addons

Istio use Prometheus to store metrics

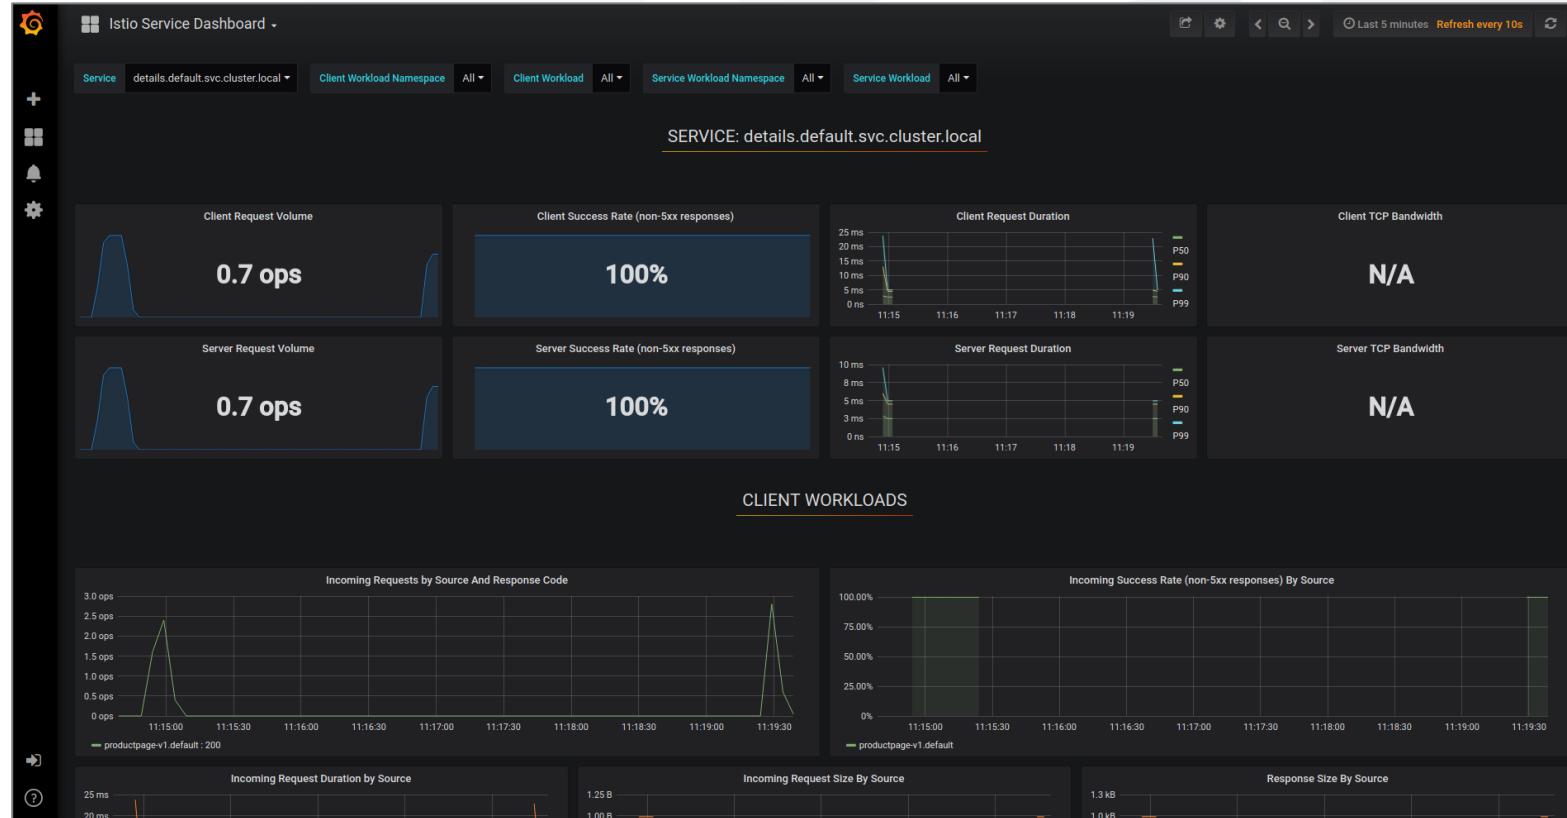
The screenshot shows the Prometheus web interface with the following details:

- Top Bar:** Prometheus, Alerts, Graph, Status ▾, Help.
- Metrics Search Bar:** istio_requests_total
- Buttons:** Execute (highlighted), - insert metric at cursor -
- Experimental UI:** Try experimental React UI
- Metrics Summary:** Load time: 151ms, Resolution: 14s, Total time series: 8
- Graph/Console Selection:** Graph (selected), Console
- Time Range:** Moment
- Table of Metrics:** Shows three rows of metrics with their values.

Element	Value
istio_requests_total{connection_security_policy="mutual_tls",destination_app="details",destination_canonical_service="details",destination_port="9080",destination_principal="spiffe://cluster.local/ns/default/sa/bookinfo-details",destination_service="details.default.svc.cluster.local",destination_service_name="details",destination_service_namespace="default",destination_version="v1",destination_workload="details-v1",destination_workload_namespace="default",instance="10.12.2.13:15090",job="envoy-stats",namespace="default",pod_name="details-v1-c5b5f496d-wsv5q",reporter="destination",request_protocol="http",response_code="200",response_flags="-",source_app="productpage",source_canonical_service="productpage",source_principal="spiffe://cluster.local/ns/default/sa/bookinfo-productpage",source_version="v1",source_workload="productpage-v1",source_workload_namespace="default"}	8
istio_requests_total{connection_security_policy="mutual_tls",destination_app="productpage",destination_canonical_service="productpage",destination_port="9080",destination_principal="spiffe://cluster.local/ns/default/sa/bookinfo-productpage",destination_service="productpage.default.svc.cluster.local",destination_service_name="productpage",destination_service_namespace="default",destination_version="v1",destination_workload="productpage-v1",destination_workload_namespace="default",instance="10.12.1.15:15090",job="envoy-stats",namespace="default",pod_name="productpage-v1-7d6cfb7dfd-rhn7l",reporter="destination",request_protocol="http",response_code="0",response_flags="DC",source_app="istio-ingressgateway",source_canonical_service="istio-ingressgateway",source_principal="spiffe://cluster.local/ns/istio-system/sa/istio-ingressgateway-service-account",source_version="unknown",source_workload="istio-ingressgateway",source_workload_namespace="istio-system"}	4
istio_requests_total{connection_security_policy="mutual_tls",destination_app="productpage",destination_canonical_service="productpage",destination_port="9080",destination_principal="spiffe://cluster.local/ns/default/sa/bookinfo-productpage",destination_service="productpage.default.svc.cluster.local",destination_service_name="productpage",destination_service_namespace="default",destination_version="v1",destination_workload="productpage-v1",destination_workload_namespace="default",instance="10.12.1.15:15090",job="envoy-stats",namespace="default",pod_name="productpage-v1-7d6cfb7dfd-rhn7l",reporter="destination",request_protocol="http",response_code="200",response_flags="-",source_app="istio-ingressgateway",source_canonical_service="istio-ingressgateway",source_principal="spiffe://cluster.local/ns/istio-system/sa/istio-ingressgateway-service-account",source_version="unknown",source_workload="istio-ingressgateway",source_workload_namespace="istio-system"}	9

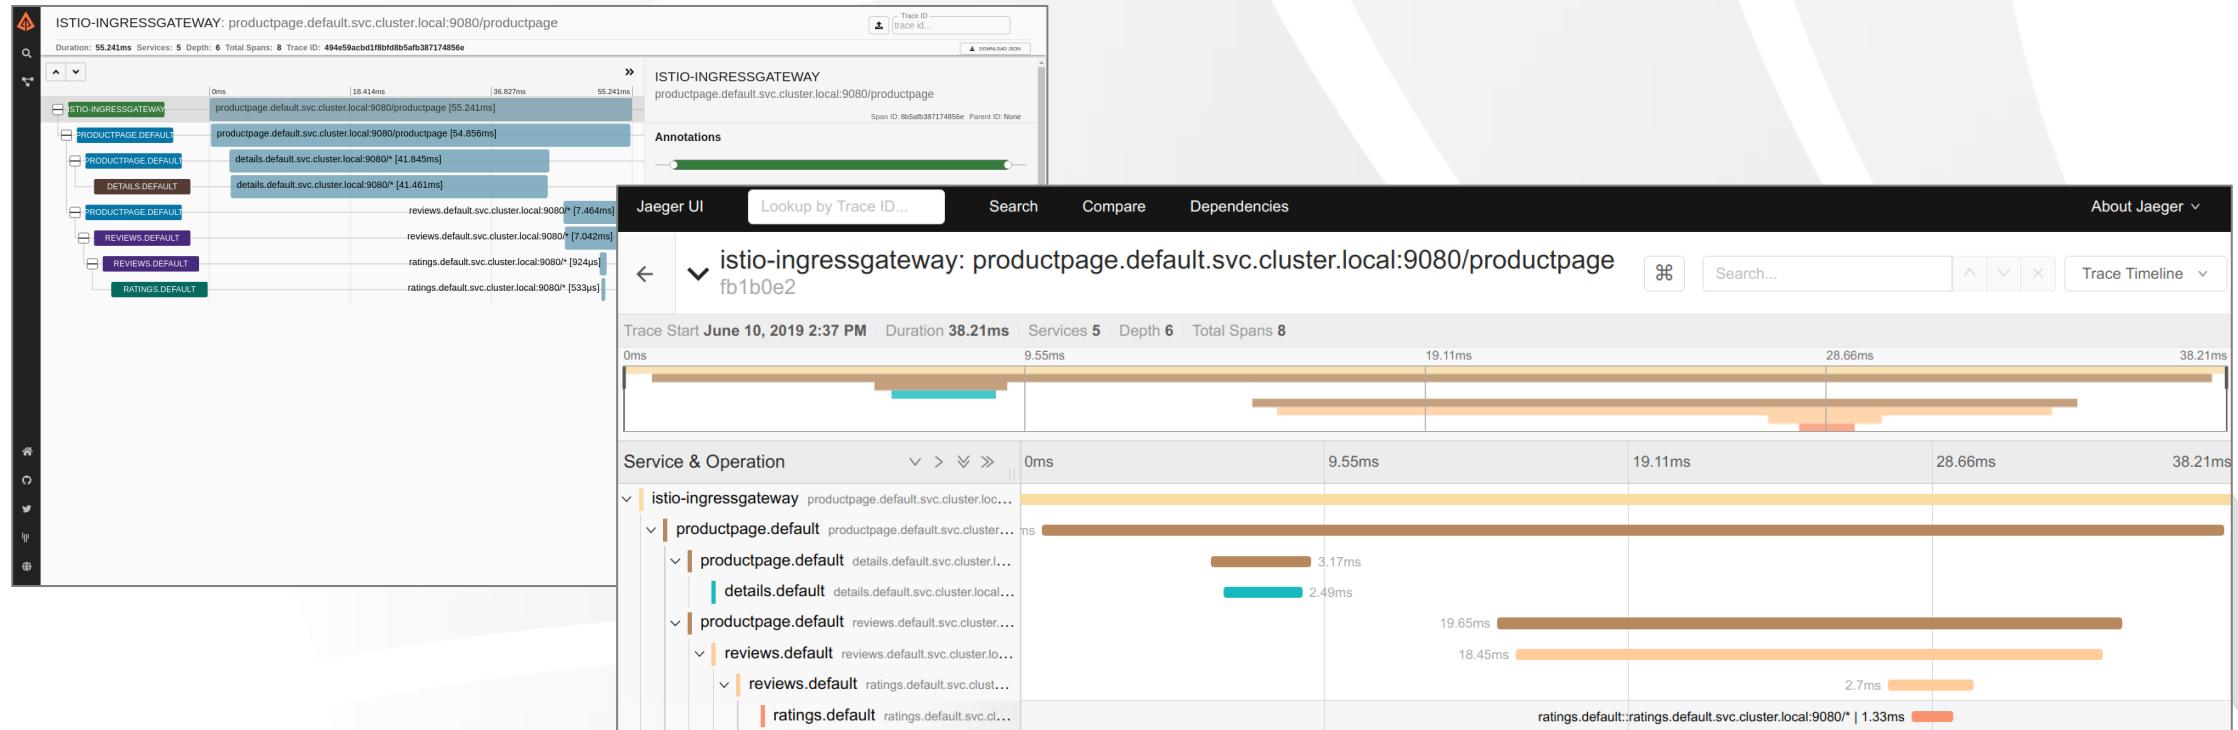
Istio Addons

Istio use Grafana for metrics visualization



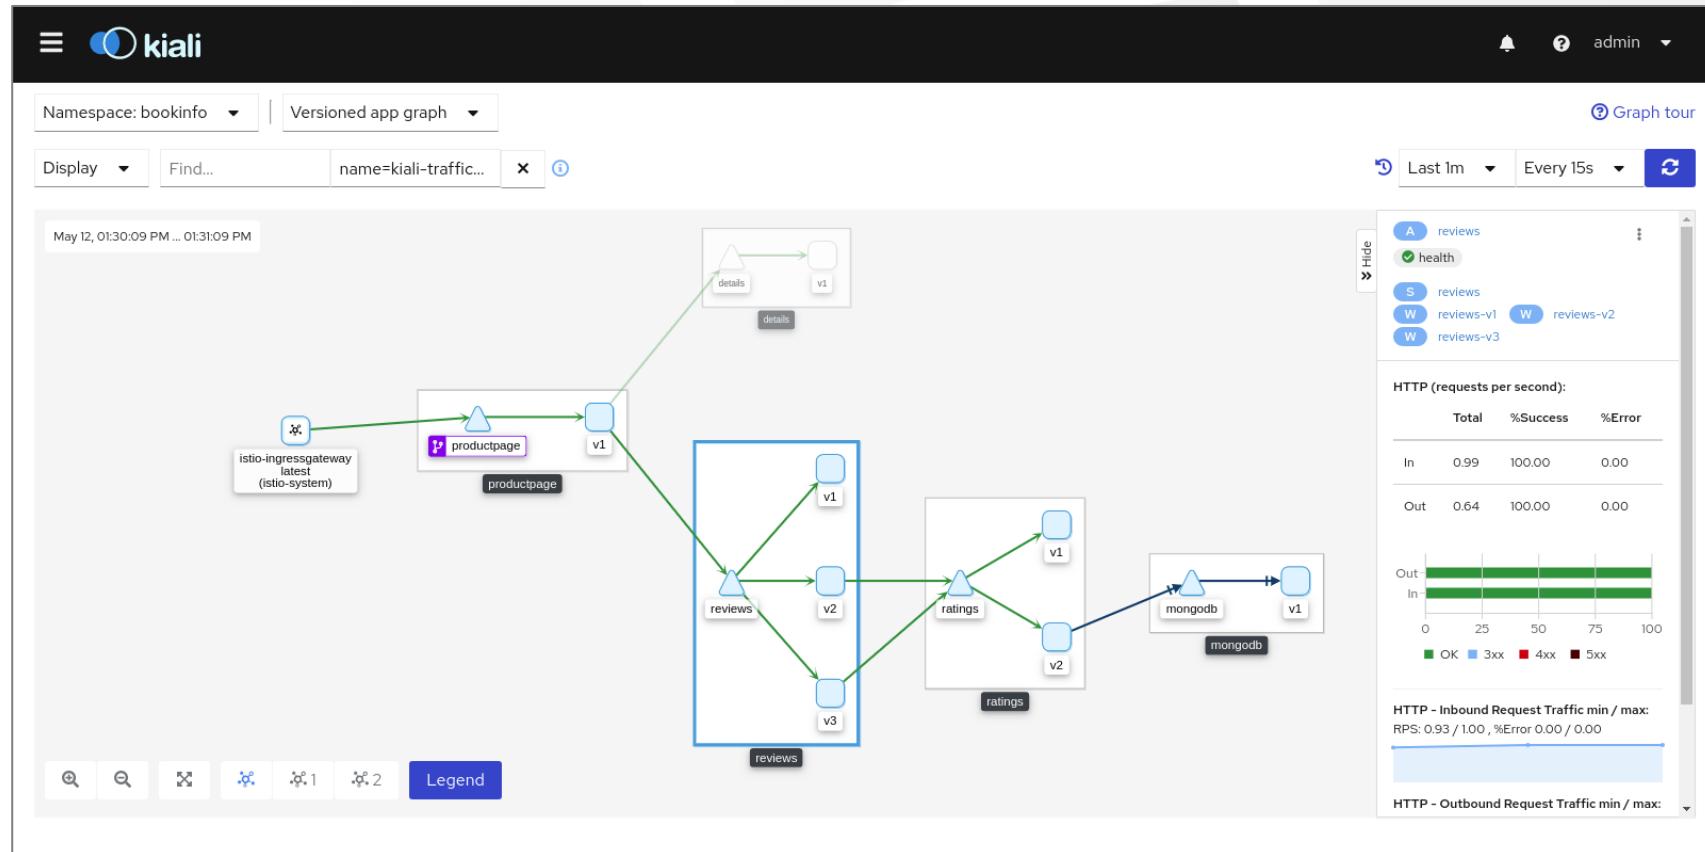
Istio Addons

Istio can use Jaeger or Zipkin for tracing visualization



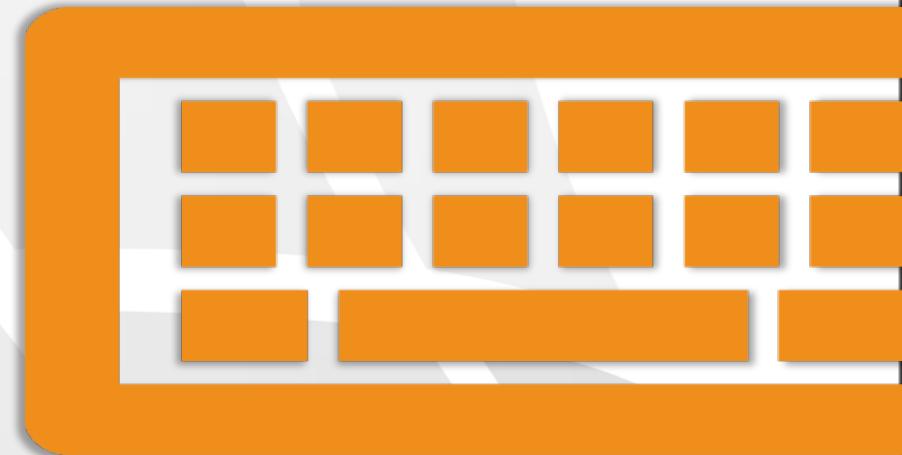
Istio Addons

Istio use Kiali for the mesh visualization



Lab 07: Installing Istio

Lab

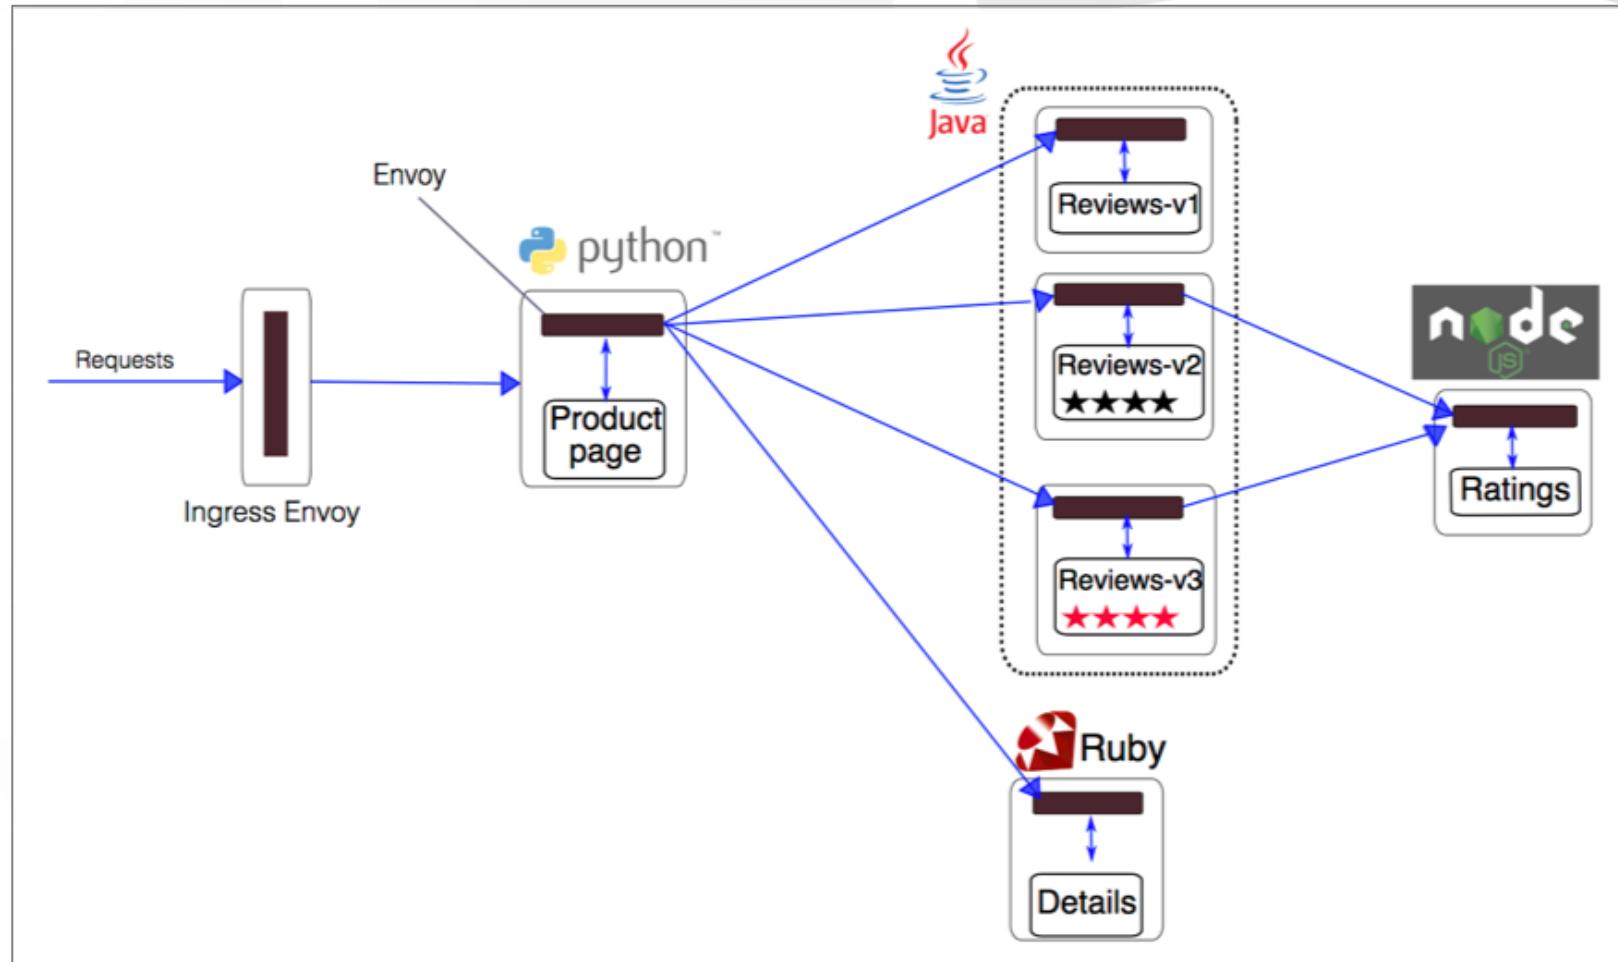


<https://github.com/leonjalfon1/sdp-servicemesh/labs/07-installing-istio.md>

BookInfo Demo App

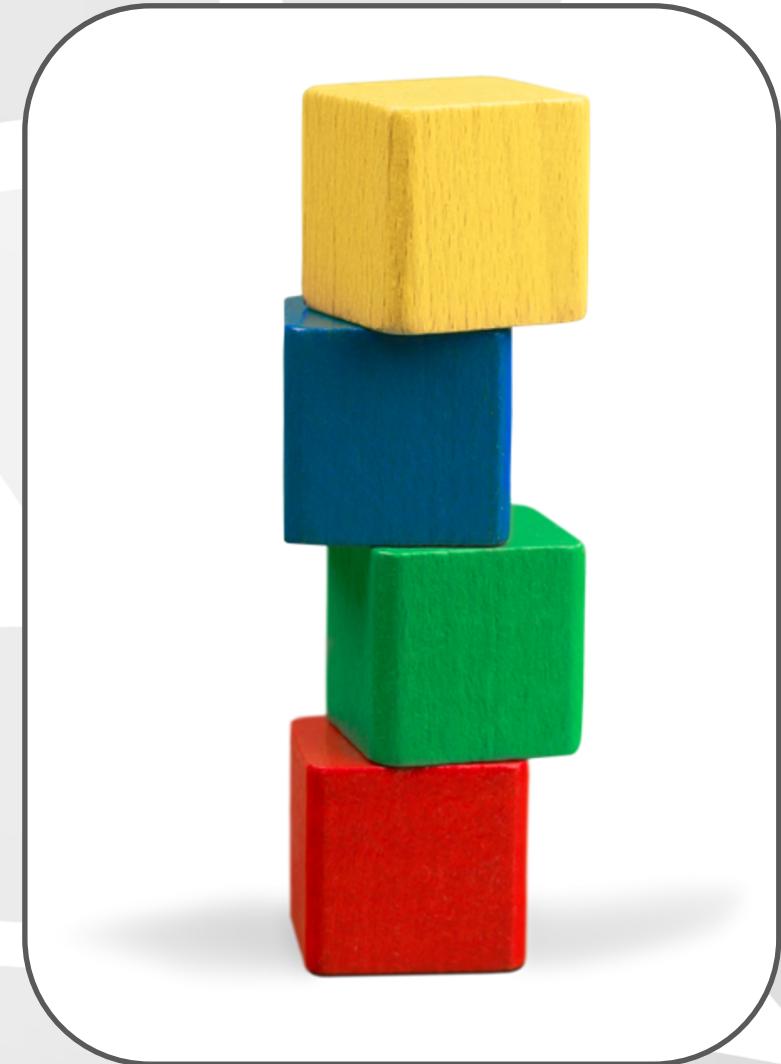
- ★ Let's use the BookInfo application to learn the following topics:
 - ★ Intelligent Routing
 - ★ Fault inject (delay/failure)
 - ★ Traffic Management
 - ★ Authorization

BookInfo – Architecture



Istio Building Blocks

- Gateway
- VirtualService
- DestinationRule
- ServiceEntry
- Sidecar
- PeerAuthentication
- AuthorizationPolicy
- ...



Gateway

- ★ You use a gateway to manage inbound and outbound traffic for your mesh, letting you specify which traffic you want to enter or leave the mesh
- ★ Gateway configurations are applied to standalone Envoy proxies that are running at the edge of the mesh, rather than sidecar Envoy proxies running alongside your service workloads
- ★ Unlike Kubernetes Ingress Controllers, Istio gateways let you use the full power and flexibility of Istio's traffic routing

Gateway

```
apiVersion: networking.istio.io/v1alpha3
kind: Gateway
metadata:
  name: bookinfo-gateway
spec:
  selector:
    istio: ingressgateway    # use istio default controller
  servers:
  - port:
      number: 80
      name: http
      protocol: HTTP
    hosts:
    - "*"
```

This gateway configuration lets HTTP traffic from any hostname into the mesh on port 80, but doesn't specify any routing for the traffic

VirtualService

- ★ A virtual service lets you configure how requests are routed to a service within an Istio service mesh
- ★ Each virtual service consists of a set of routing rules that are evaluated in order
- ★ Istio matches each given request to the virtual service to a specific real destination within the mesh
- ★ Without virtual services, Envoy distributes traffic using round-robin load balancing between all service instances

VirtualService

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - "*"
  gateways:
    - bookinfo-gateway
  http:
    - route:
        - destination:
            host: reviews
            subset: v1
```

This VirtualService route all the traffic coming from the “bookinfo-gateway” to the destination “reviews” with subset “v1”

DestinationRule

- ★ You can think of virtual services as how you route your traffic to a given destination, and then you use destination rules to configure what happens to traffic for that destination
- ★ Destination rules are applied after virtual service routing rules are evaluated, so they apply to the traffic's "real" destination
- ★ Each subset is defined based on one or more labels
- ★ By default, Istio uses a round-robin load balancing policy but you can specify another algorithm in destination rules

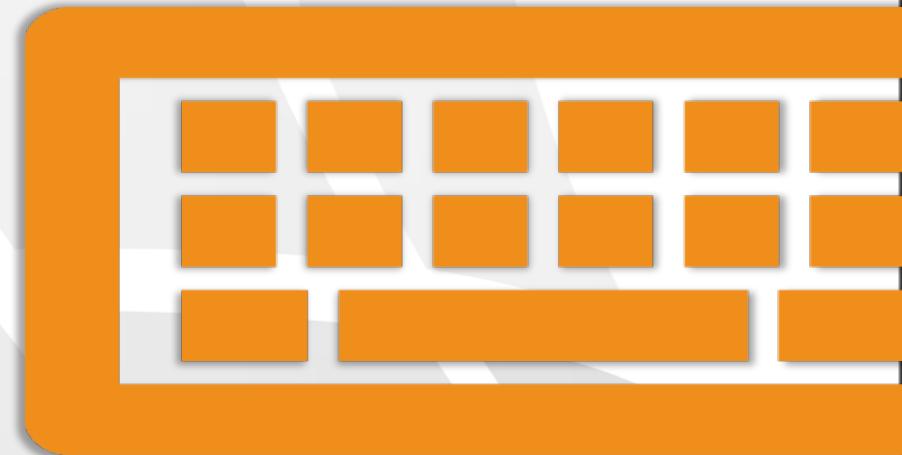
DestinationRule

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
    - name: v3
      labels:
        version: v3
```

This destination rule define three subsets for the reviews service. For example, the destination “v1” will route the traffic to a Kubernetes service using the label “version=v1”

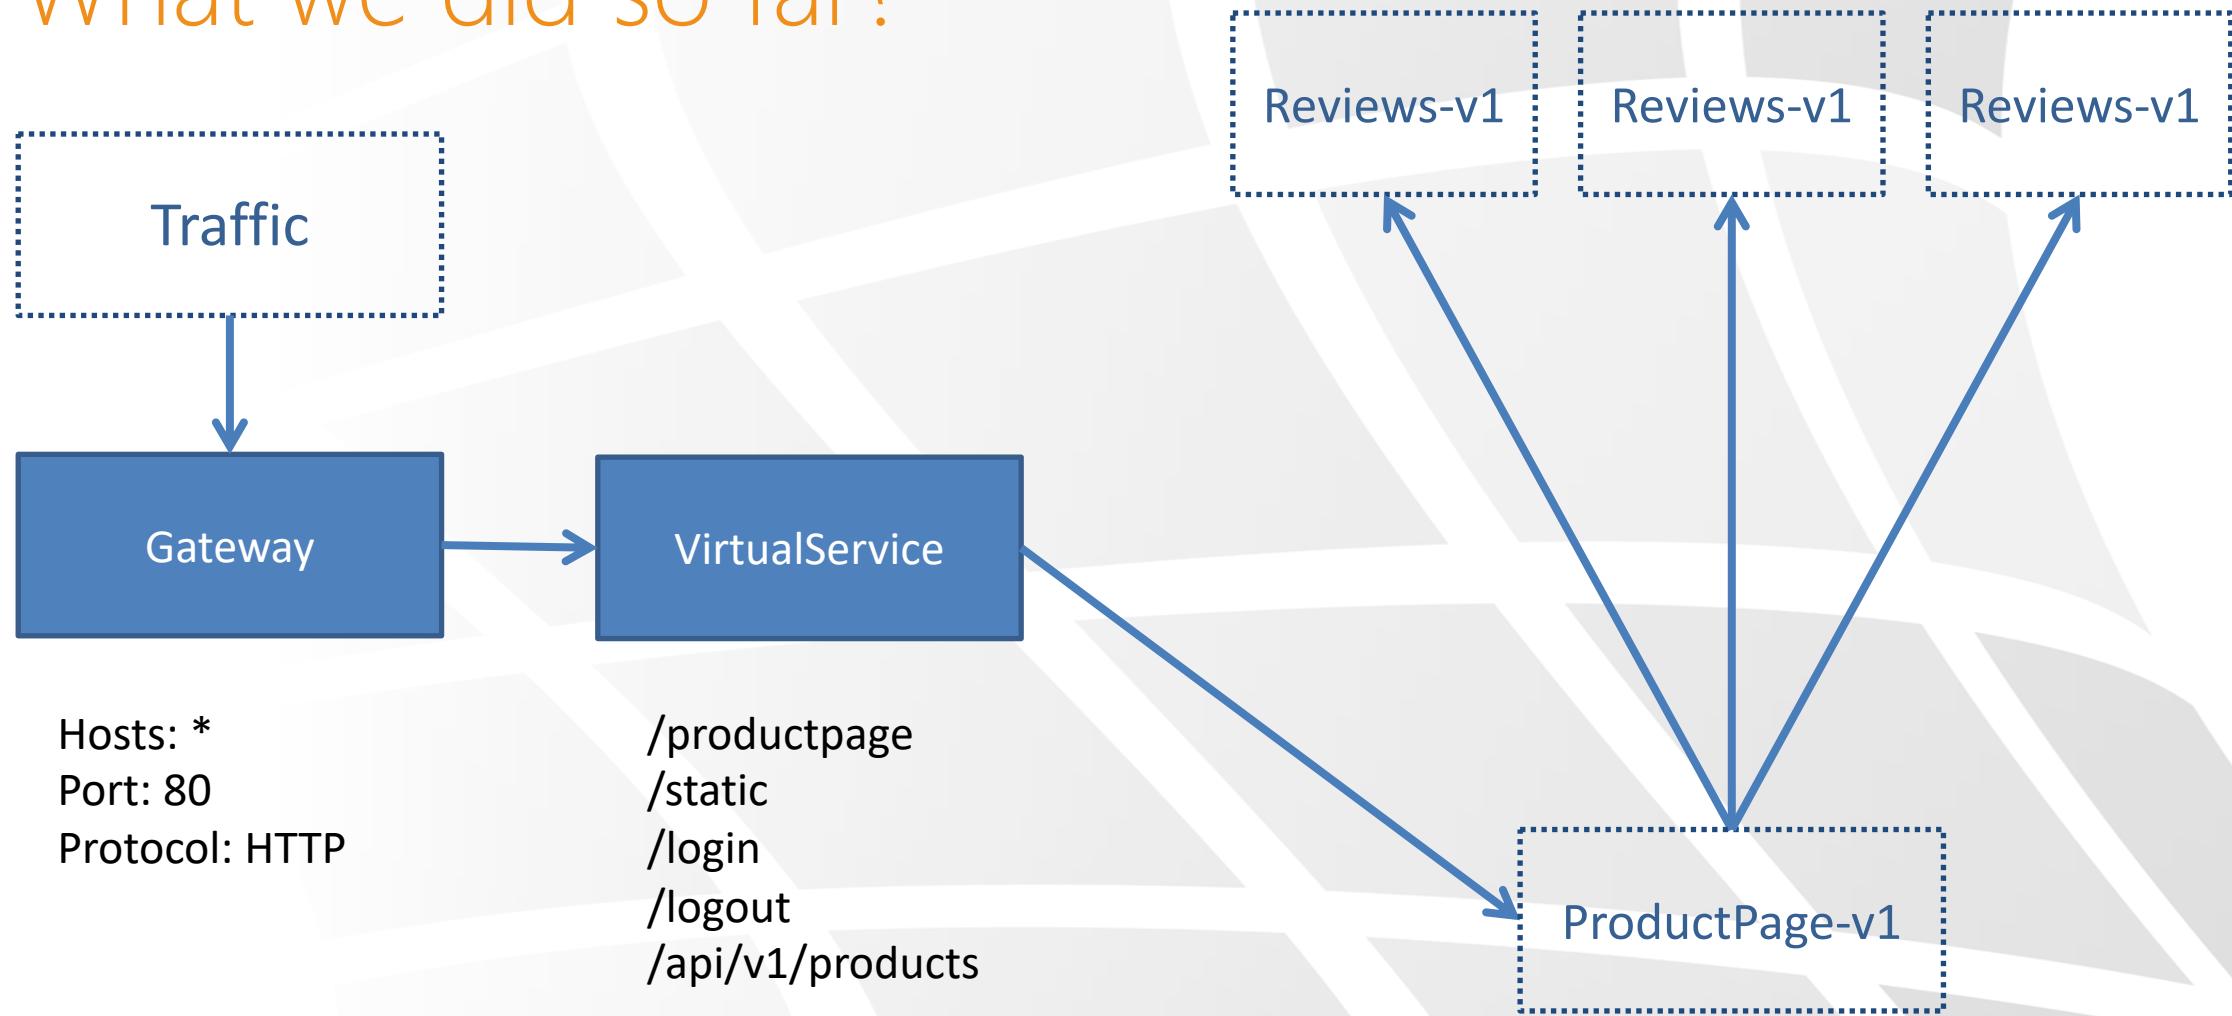
Lab 08: Exploring Istio

Lab

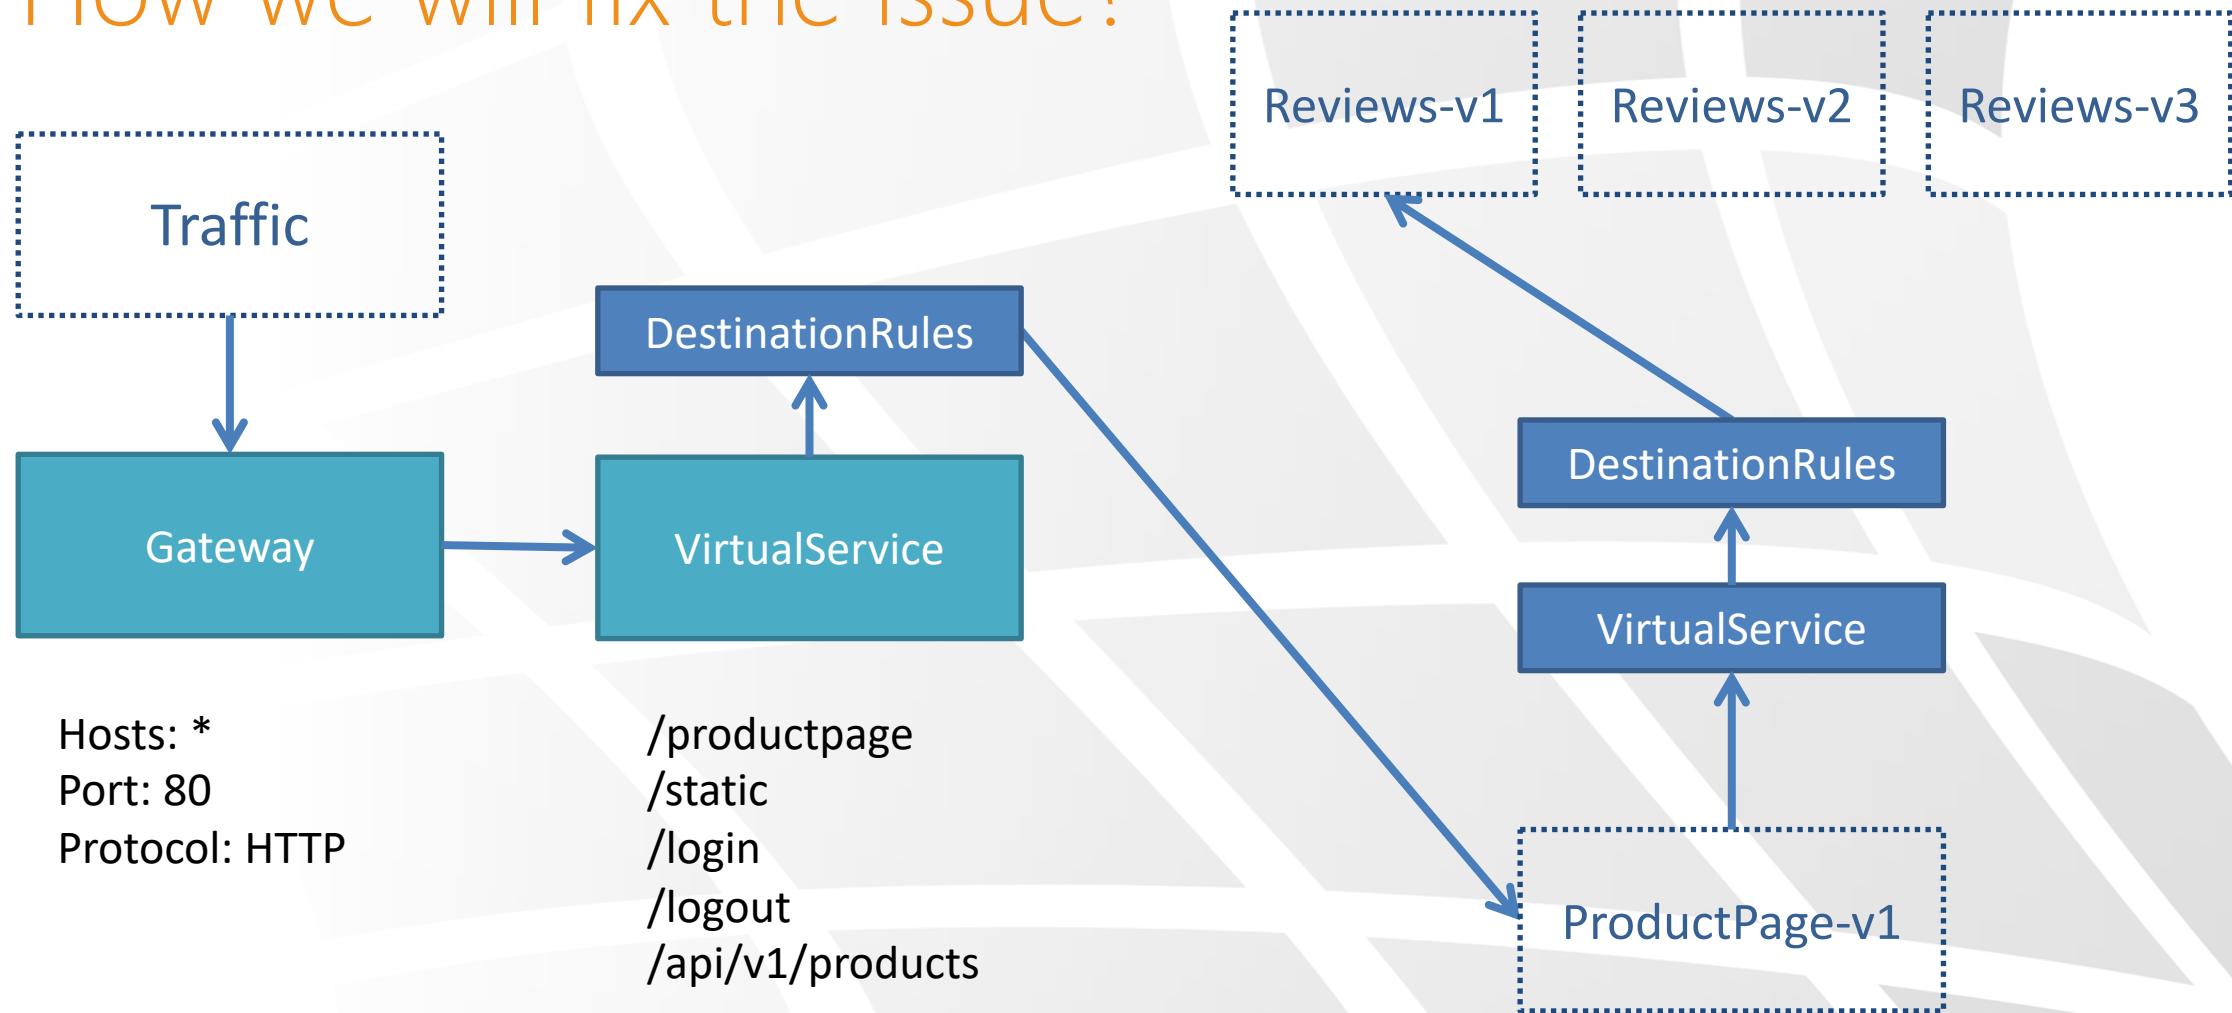


<https://github.com/leonjalfon1/sdp-servicemesh/labs/08-exploring-istio.md>

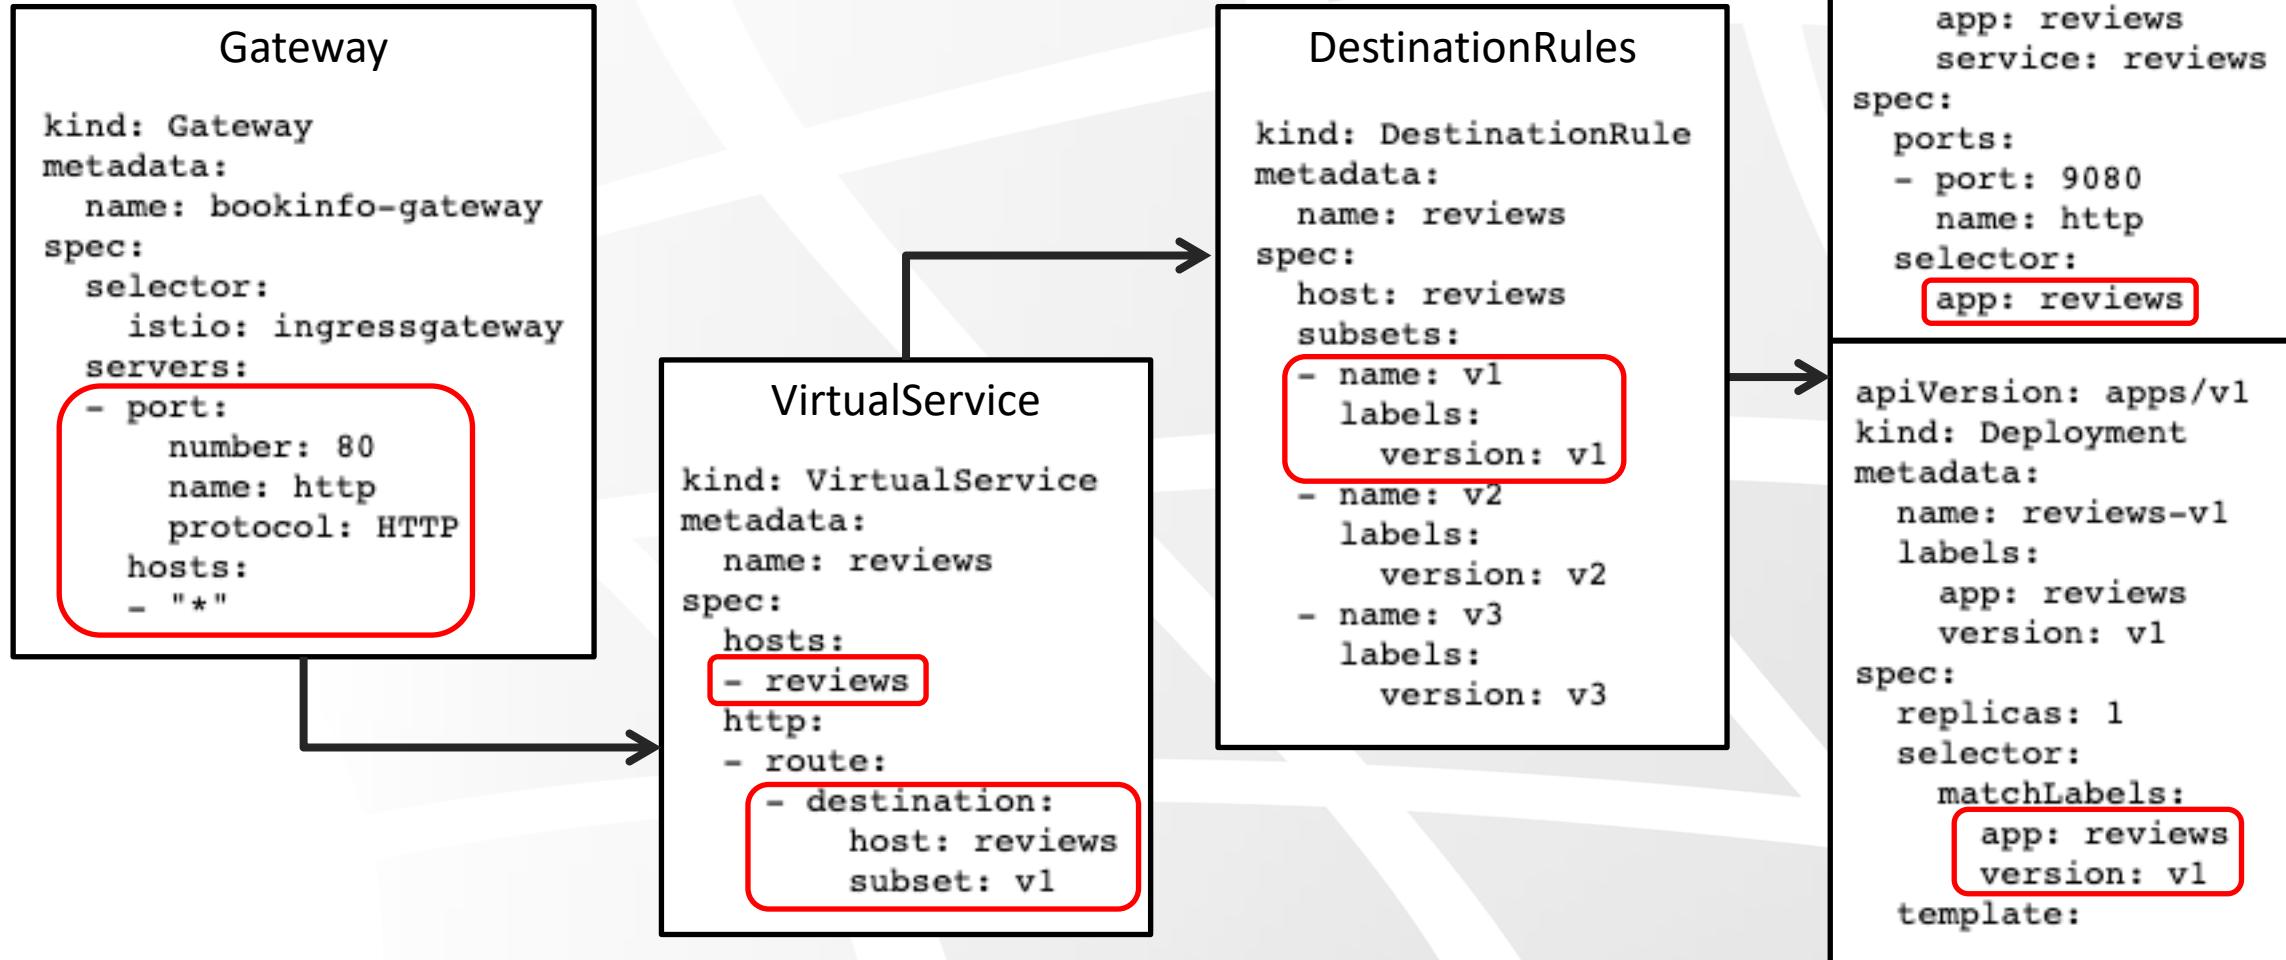
What we did so far?



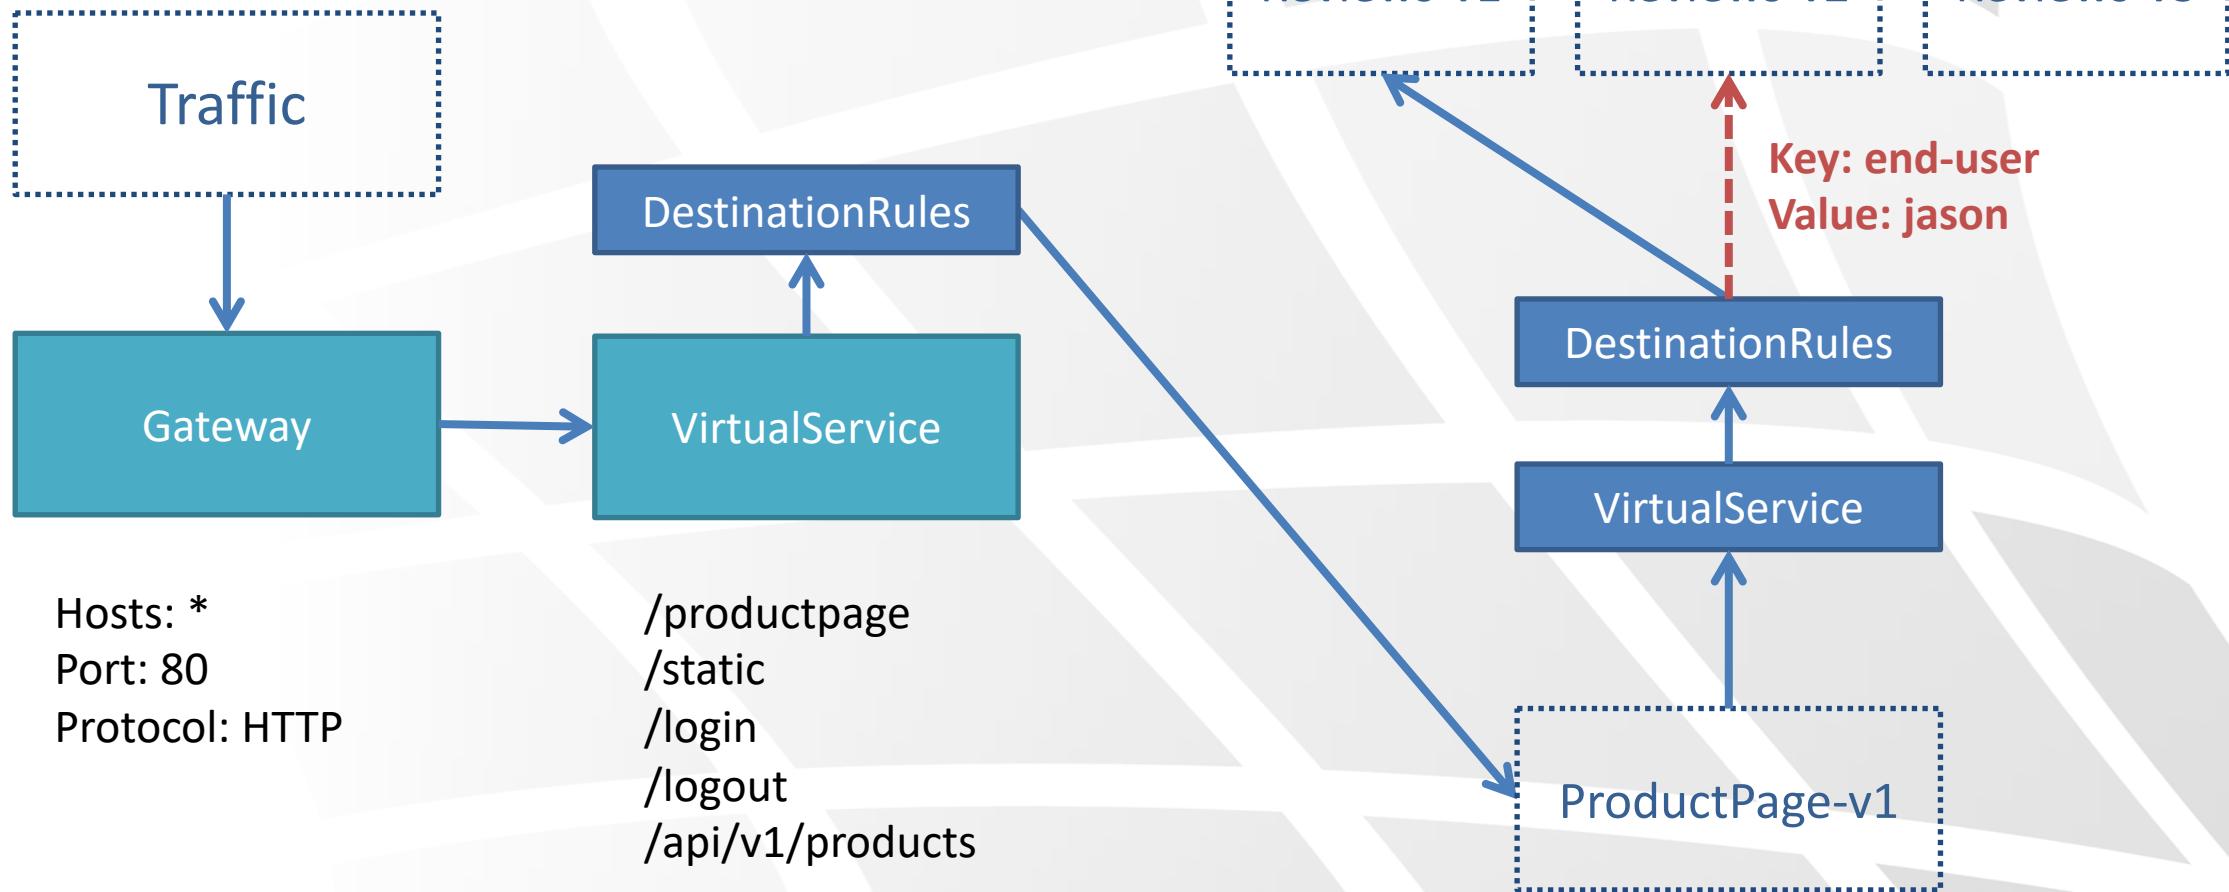
How we will fix the issue?



Let's take a look into the puzzle



One more thing...

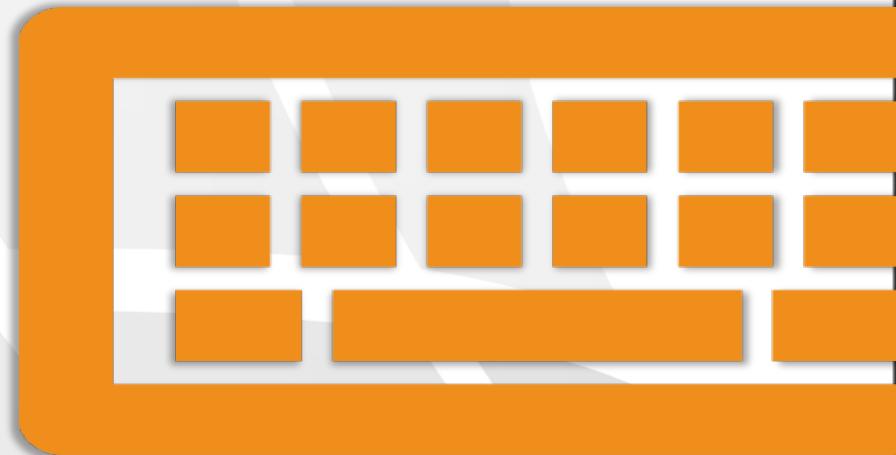


VirtualService – Match Rules

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - match:
        - headers:
            end-user:
              exact: jason
        route:
          - destination:
              host: reviews
              subset: v2
    - route:
        - destination:
            host: reviews
            subset: v1
```

Lab 09: Intelligent Routing with Istio

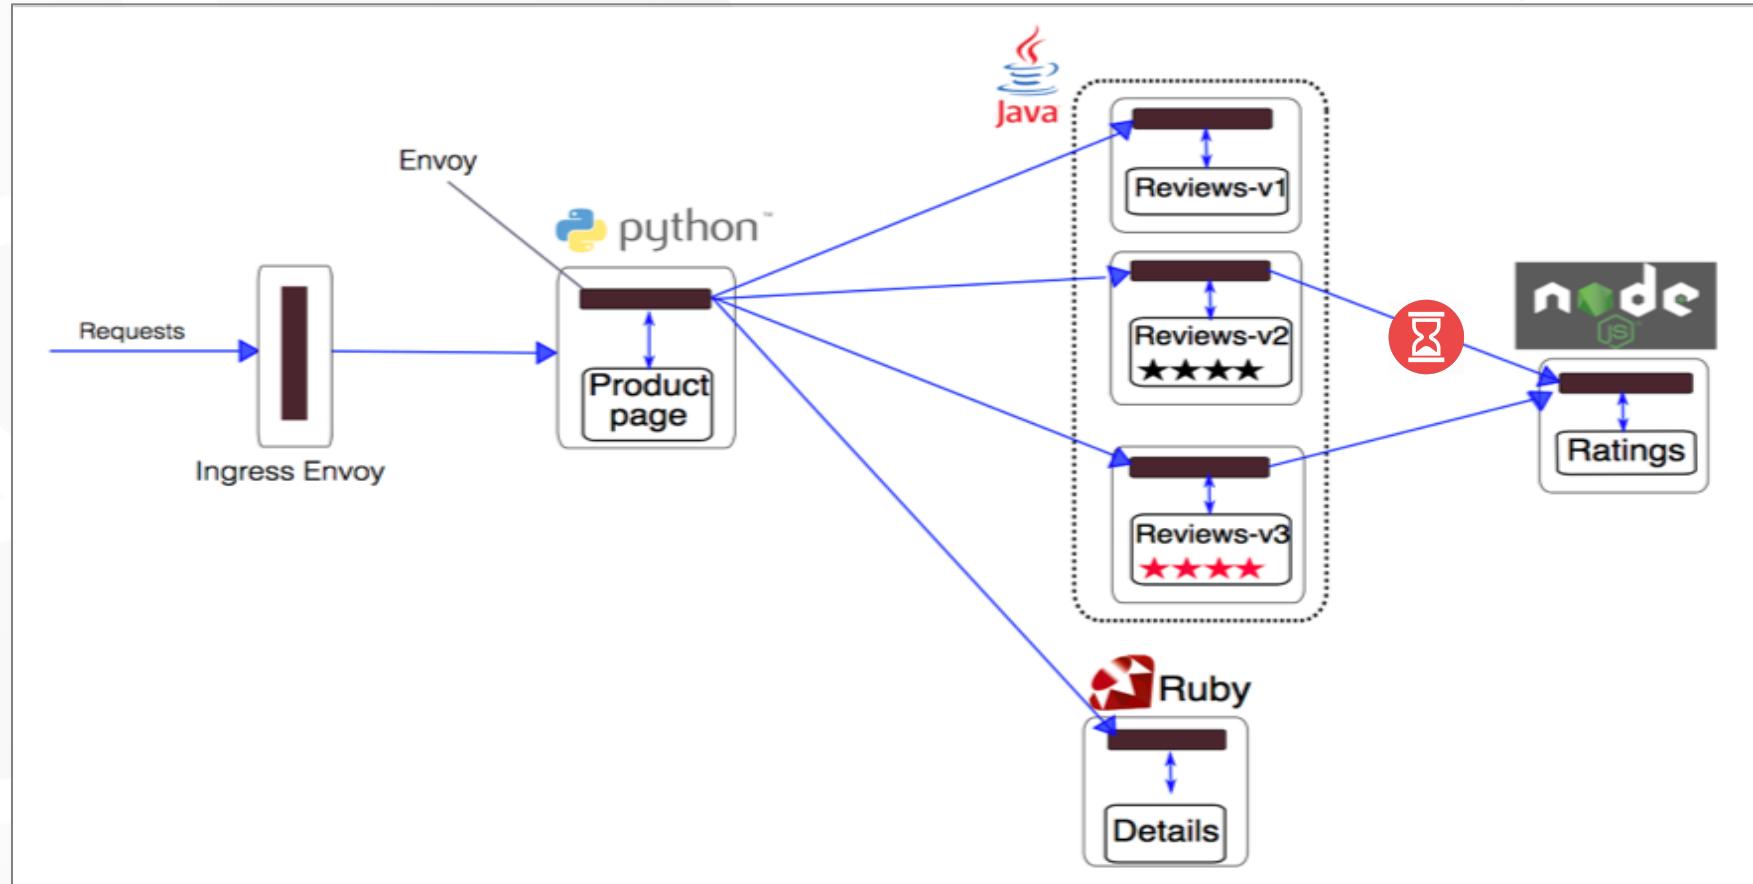
Lab



<https://github.com/leonjalfon1/sdp-servicemesh/labs/09-intelligent-routing-istio.md>

Delay Injection

❖ Delay

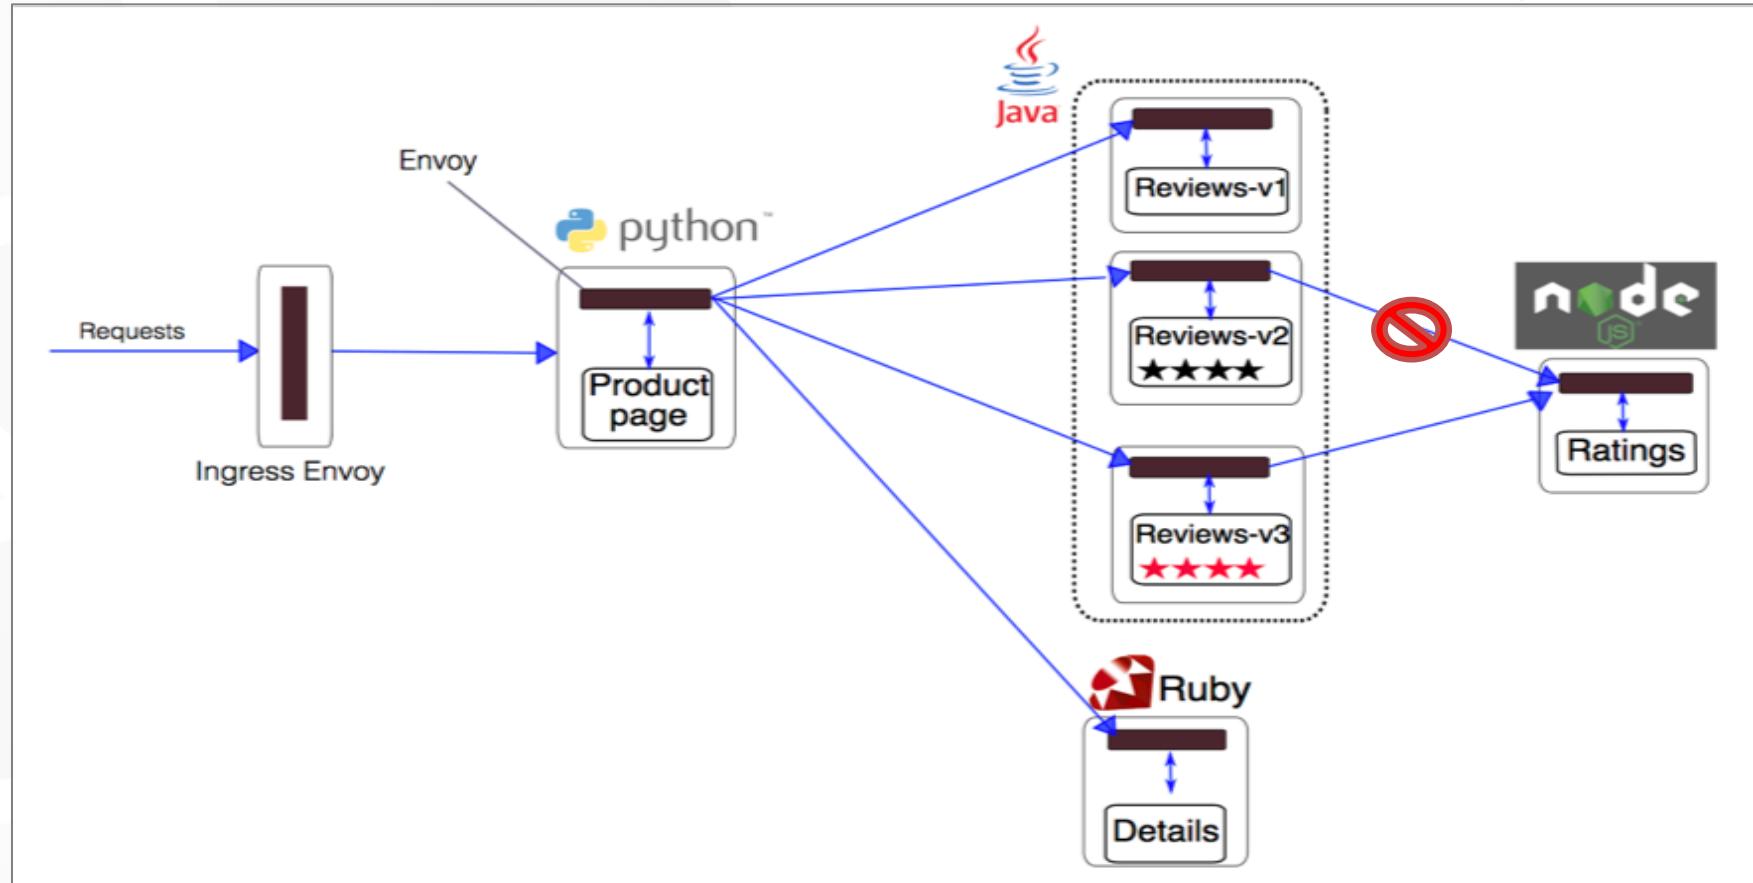


VirtualService – Delay Injection

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    - headers:
      end-user:
        exact: jason
    fault:
      delay:
        percentage:
          value: 100.0
        fixedDelay: 7s
    route:
    - destination:
        host: ratings
        subset: v1
    - route:
      - destination:
          host: ratings
          subset: v1
```

Error Injection

❖ Error

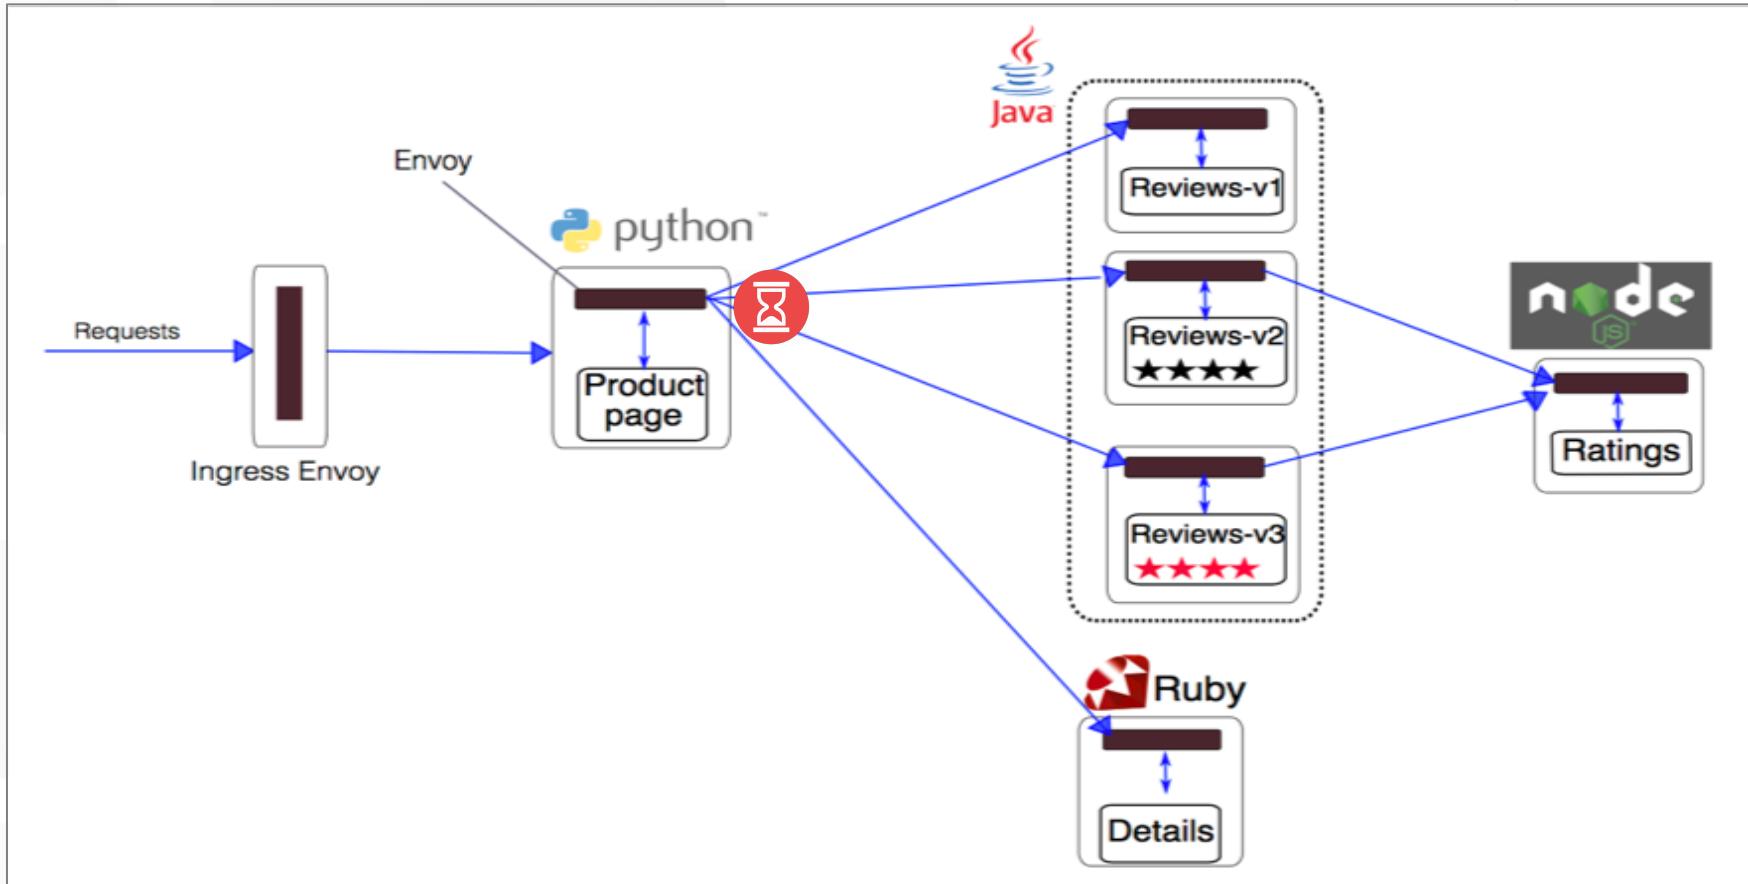


VirtualService – Error Injection

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    - headers:
        end-user:
          exact: jason
    fault:
      abort:
        percentage:
          value: 100.0
        httpStatus: 500
    route:
    - destination:
        host: ratings
        subset: v1
    - route:
        - destination:
            host: ratings
            subset: v1
```

Requests Timeout

⌚ Timeout

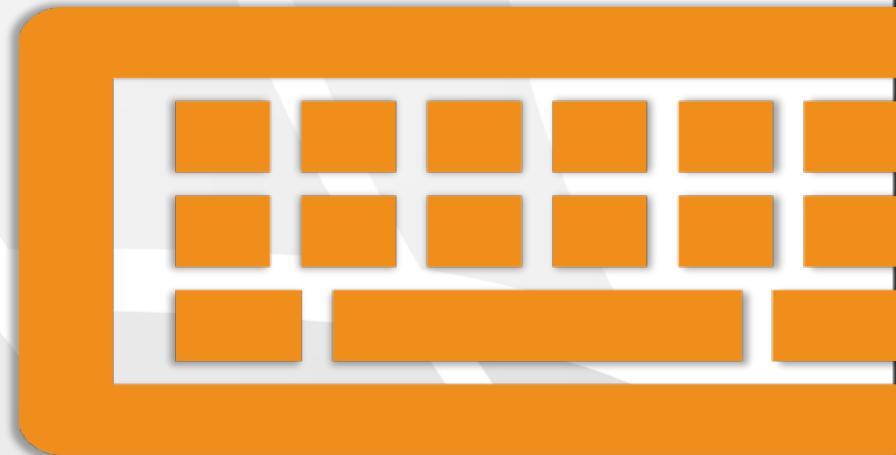


VirtualService – Requests Timeout

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - route:
      - destination:
          host: reviews
          subset: v2
        timeout: 0.5s
```

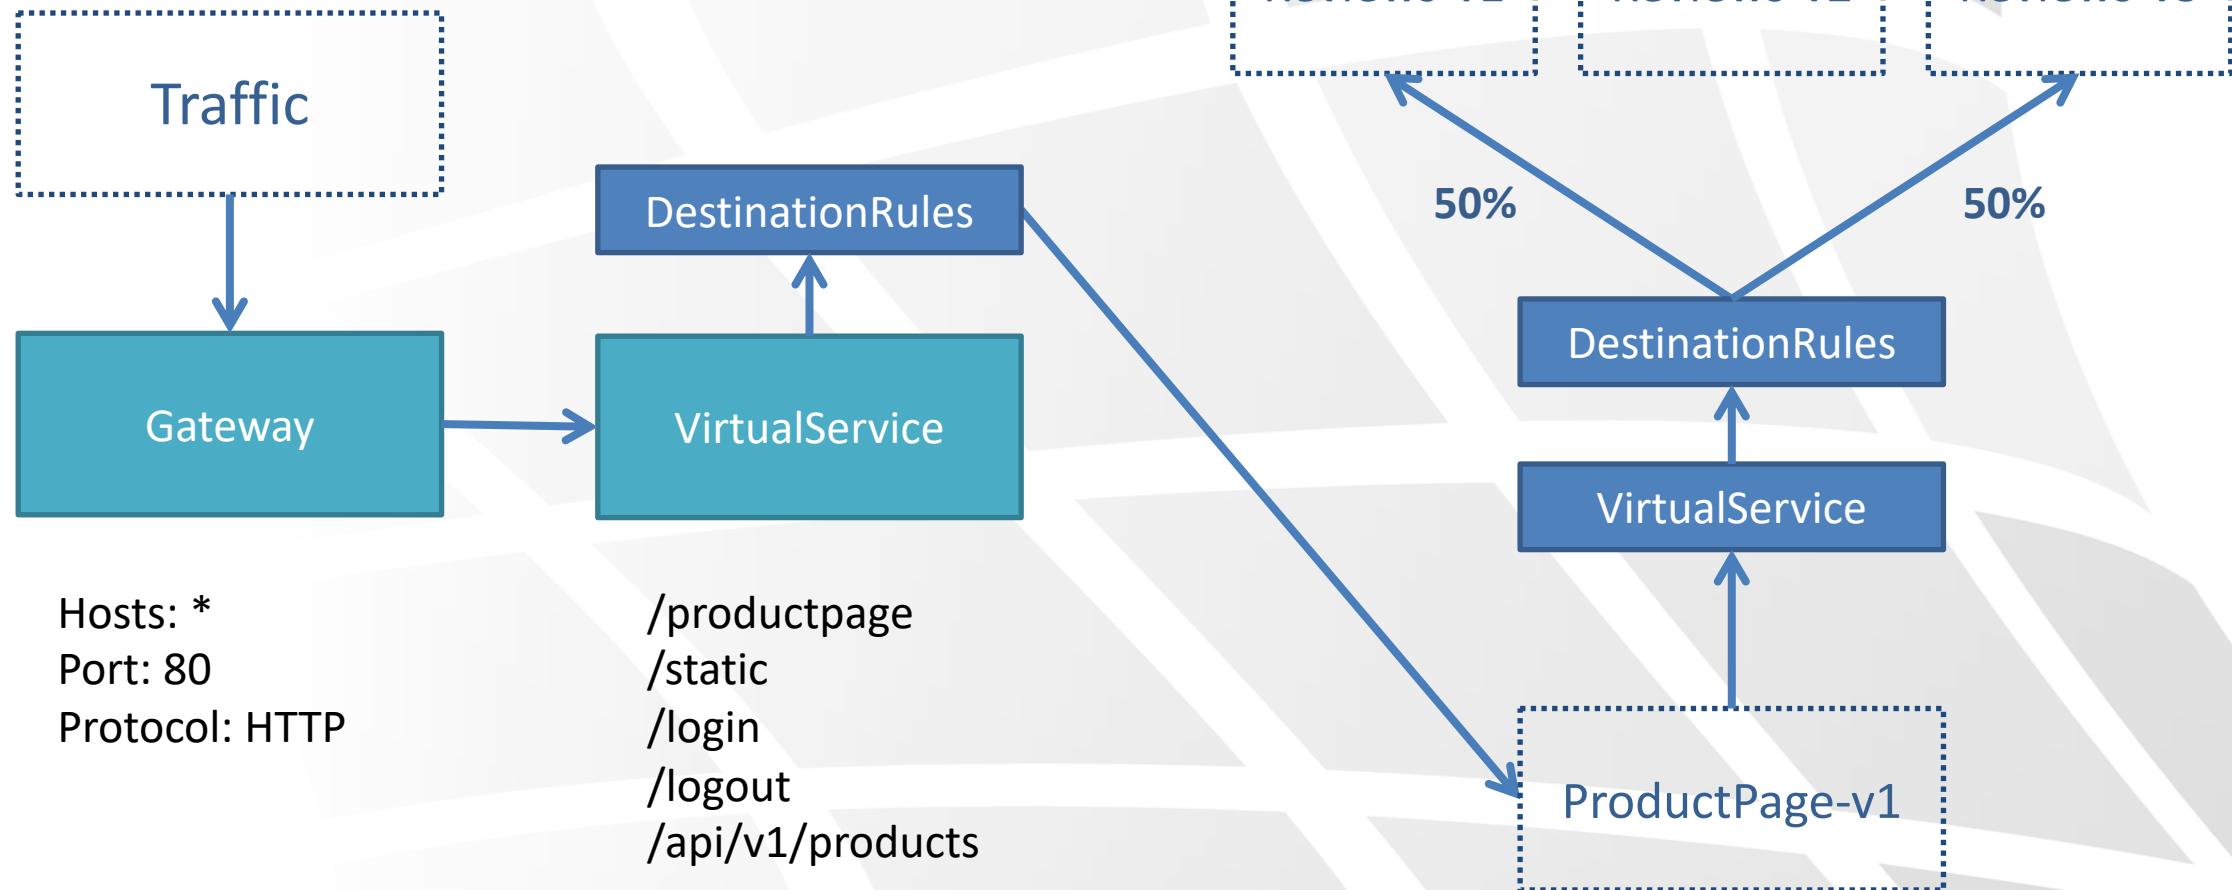
Lab 10: Fault Injection with Istio

Lab



<https://github.com/leonjalfon1/sdp-servicemesh/labs/10-fault-injection-istio.md>

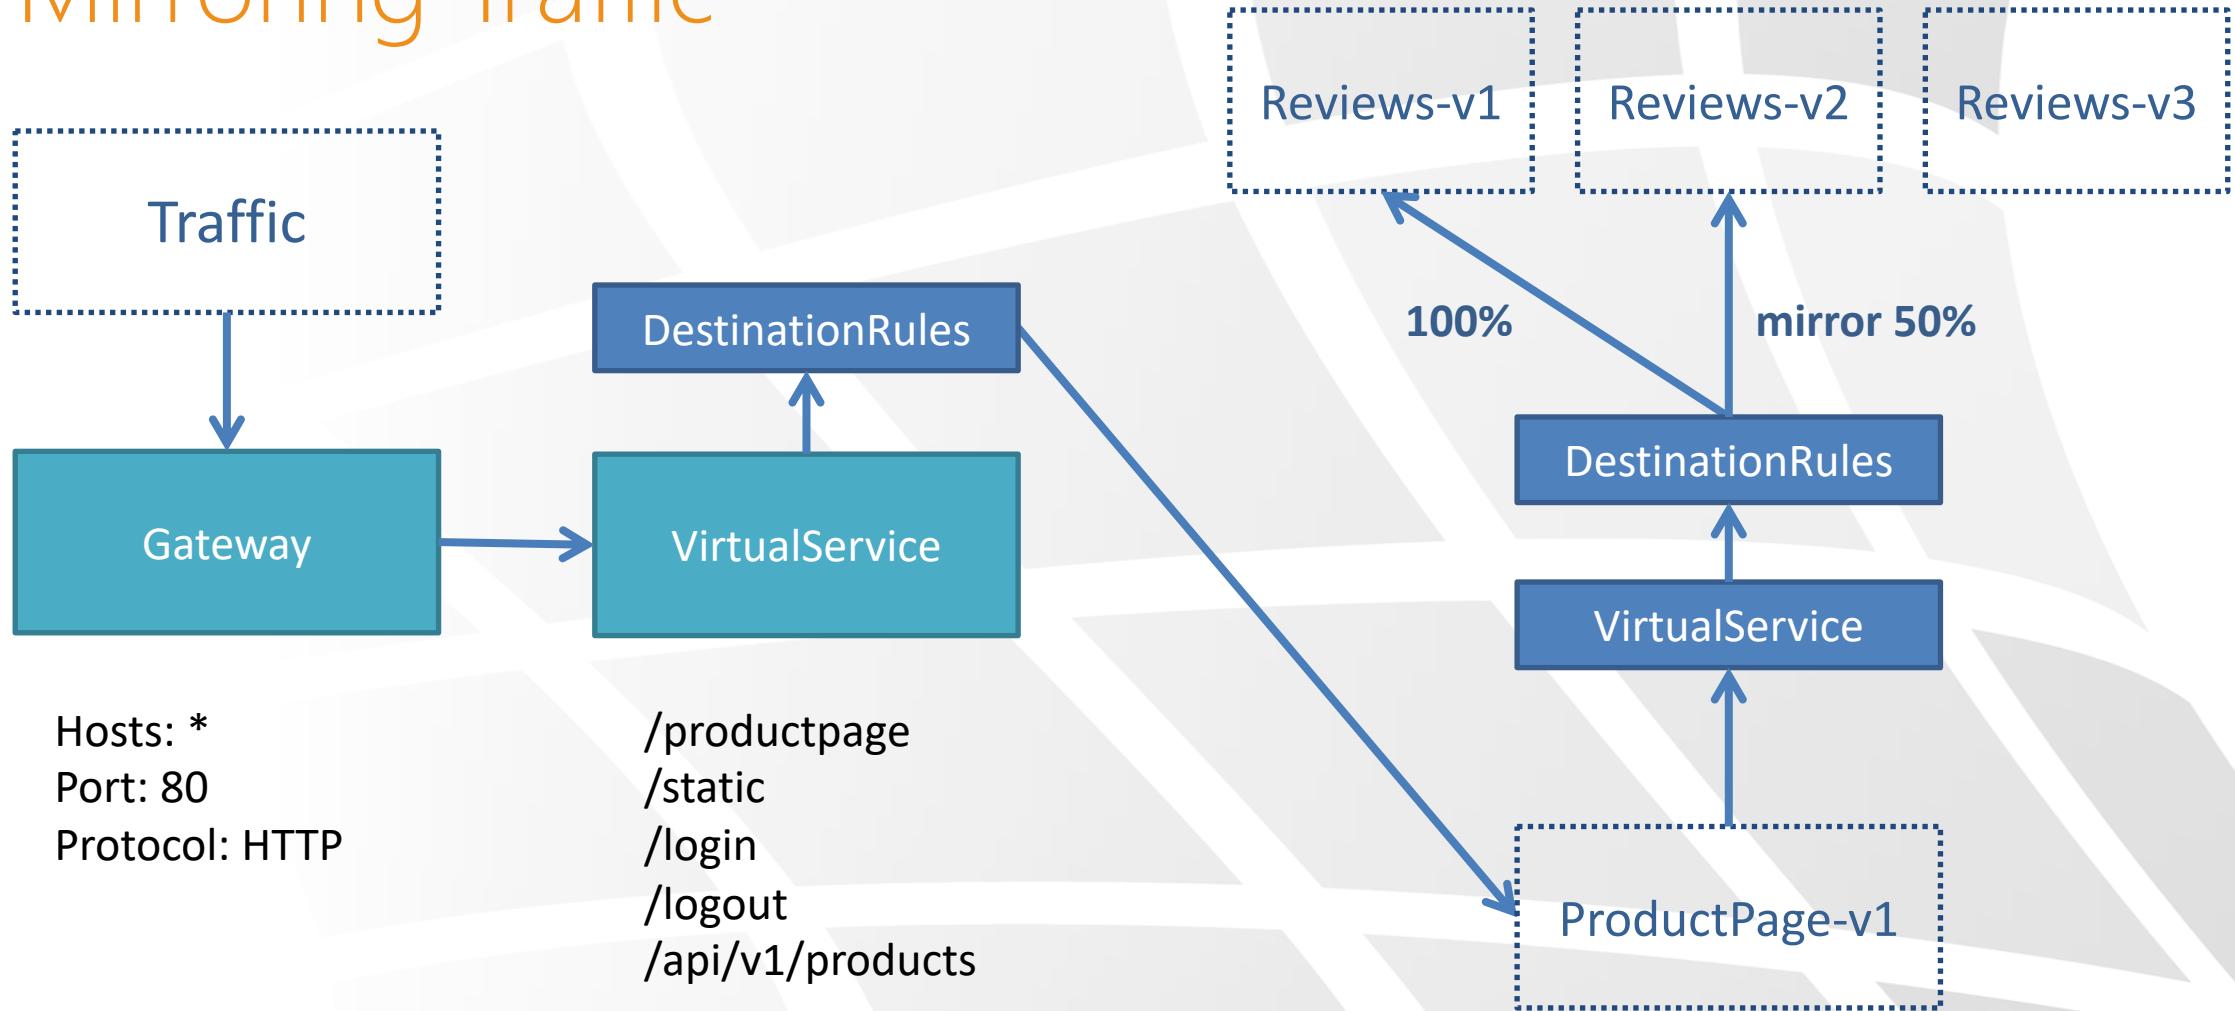
Canary Deployments



VirtualService – Traffic Shifting

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - route:
        - destination:
            host: reviews
            subset: v1
            weight: 50
        - destination:
            host: reviews
            subset: v3
            weight: 50
```

Mirroring Traffic

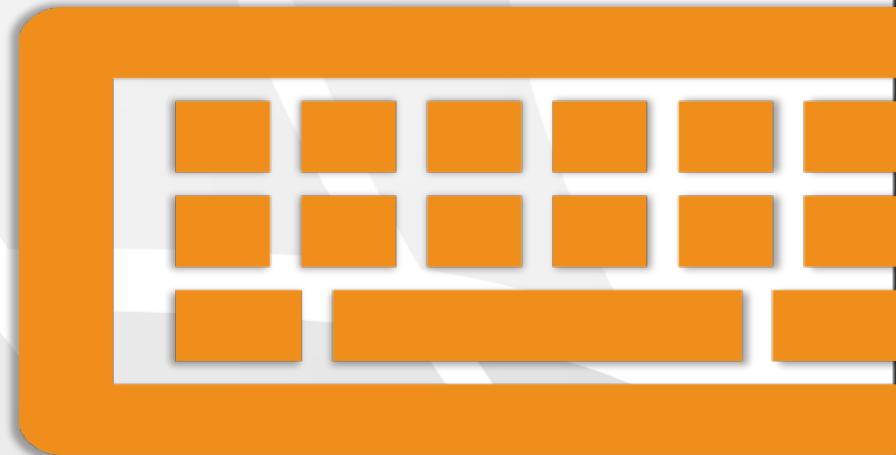


VirtualService – Mirroring Traffic

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - route:
      - destination:
          host: reviews
          subset: v1
          weight: 100
        mirror:
          host: reviews
          subset: v2
          mirror_percent: 50
```

Lab 11: Traffic Management with Istio

Lab



<https://github.com/leonjalfon1/sdp-servicemesh/labs/11-traffic-management-istio.md>

Istio Authorization Benefits

- ★ Workload-to-workload and end-user-to-workload authorization.
- ★ A Simple API: it includes a single AuthorizationPolicy CRD, which is easy to use and maintain.
- ★ Flexible semantics: operators can define custom conditions on Istio attributes, and use DENY and ALLOW actions.
- ★ High performance: Istio authorization is enforced natively on Envoy.
- ★ High compatibility: supports gRPC, HTTP, HTTPS and HTTP2 natively, as well as any plain TCP protocols.

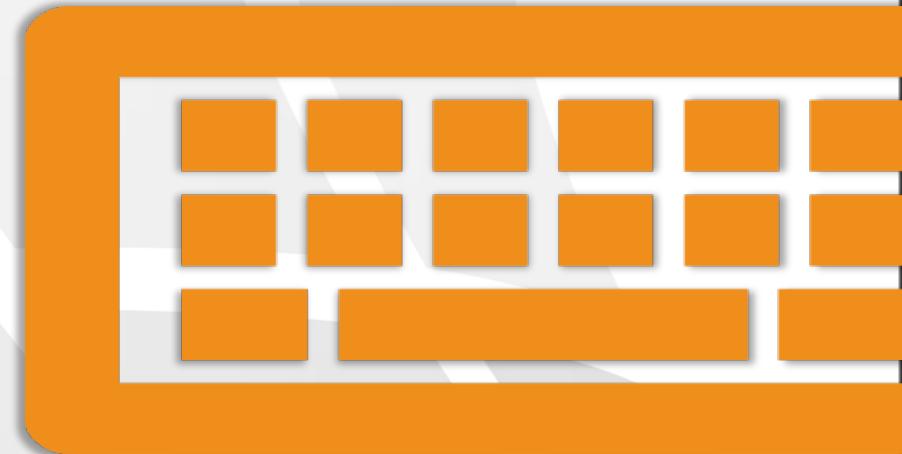
Authorization Policy (CRD)

- ★ The selector field specifies the target of the policy
- ★ The action field specifies whether to allow or deny the request
- ★ The rules specify when to trigger the action

```
apiVersion: "security.istio.io/v1beta1"
kind: "AuthorizationPolicy"
metadata:
  name: "details-viewer"
  namespace: default
spec:
  selector:
    matchLabels:
      app: details
  rules:
    - from:
        - source:
            principals: ["cluster.local/ns/default/sa/bookinfo-productpage"]
        to:
        - operation:
            methods: ["GET"]
```

Lab 12: Setup Istio Authorization

Lab



<https://github.com/leonjalfon1/sdp-servicemesh/labs/12-setup-istio-authorization.md>

IT FEELS LIKE THE END BUT



Spongebob Squarepants vector trace by kssael. ©Nickelodeon

IT'S ONLY THE BEGINNING

The Service Mesh Interface (SMI)

- ★ A standard interface for service meshes on Kubernetes
- ★ A basic feature set for the most common service mesh use cases
- ★ Flexibility to support new service mesh capabilities over time
- ★ Space for the ecosystem to innovate with service mesh technology

What SMI covers?

- ★ **Traffic policy** – apply policies like identity and transport encryption across services
- ★ **Traffic telemetry** – capture key metrics like error rate and latency between services
- ★ **Traffic management** – shift traffic between different services

Service Mesh Interface Ecosystem



SMI API's

- ★ **Traffic Access Control:** Apply policies like identity and transport encryption across services
- ★ **Traffic Split:** Shift traffic between different services.
- ★ **Traffic Specs:** Describe traffic on a per-protocol basis.
- ★ **Traffic Metrics:** Capture key metrics like error rate and latency between services.

TrafficSplit Example

Canary example:

```
kind: TrafficSplit
metadata:
  name: canary
spec:
  # The root service that clients use to connect to the destination application.
  service: website
  # Services inside the namespace with their own selectors, endpoints and configuration.
  backends:
  - service: website-v1
    weight: 90
  - service: website-v2
    weight: 10
```

The above configuration will route 10% of the `website` incoming traffic to the `website-v1` service and 90% to `website-v2` service.

Ok, But Why SMI?

- ❖ The goal of the SMI API is to provide a common, portable set of service mesh APIs which a Kubernetes user can use in a provider agnostic manner.
- ❖ In this way people can define applications that use service mesh technology without tightly binding to any specific implementation.

WE DID IT!

WOOHOO

What We Learned Today?

- ★ Service Mesh (with Linkerd and Istio), including:
 - ★ Installation
 - ★ Observability
 - ★ Debugging
 - ★ Fault Injection
 - ★ Intelligent Routing
 - ★ Traffic Management
 - ★ Authorization for HTTP traffic
- ★ Service Mesh Interface

So, How to Choose a Service Mesh?

	Linkerd	Istio
Advantages	Linkerd 2 is designed to be non-invasive and is optimized for performance and usability. Therefore, it requires little time to adopt.	Istio can be adapted and extended like no other Mesh. Its many features are available for Kubernetes and other platforms.
Drawbacks	Linkerd 2 is deeply integrated with Kubernetes and cannot be expanded. Since Linkerd 2 does not rely on a third-party proxy, it cannot be extended easily.	Istio's flexibility can be overwhelming for teams who don't have the capacity for more complex technology. Also, Istio takes control of the ingress controller.

[SIMPLE]

[POWERFUL]

Full comparission available here: <https://servicemesh.es/>

A cartoon illustration of a man with a large head and a determined expression, wearing a green t-shirt and blue shorts, running towards the right. He is shown from the waist up, with his arms pumping and legs moving. The background is a simple blue gradient.

**PEOPLE RUNNING
ISTIO ON K8S**

**PEOPLE RUNNING
LINKERD ON K8S**



SELA|DEVELOPER|PRACTICE
December 6-8, 2020

Thank you for your attention!

Service Mesh Fundamentals

Leon Jalfon

leonj@sela.co.il